

## Today's lecture

1. Transactions
2. Views
3. Access Rights in SQL

# Transactions

- A logical unit of work consisting of one or more SQL statements
- Atomic transaction
  - **Fully executed** or
  - **Rolled back** as if it never occurred
- Isolation from concurrent transactions
  - Changes made by a transaction are not visible to other concurrently executing transactions until the transaction completes
- Transaction model based on two SQL statements:
  - **COMMIT**
  - **ROLLBACK**
- Transactions begin implicitly
  - Ended by commit work or rollback work
- Default on most databases: each SQL statement commits automatically
- Can turn off auto commit for a session (e.g. using API)

# VIEWS

- Definition
- View Creation and Destruction
- Updating Views
- Types of Views

# Views

- One database often supports multiple applications
  - Slightly different pictures of the world.
- **Views** help accommodate this variation without storing redundant data.

## Views

- In some cases, it is not desirable for all users to see the entire logical model (that is, all the actual relations stored in the database.)

### Example:

```
Employee(ssn, name, department, project, salary)
```

Consider a person who needs to know the name and project of employees in the 'Development' department, but not the salary. This person should see a relation described, in SQL, by


```
SELECT name, project
FROM   Employee
WHERE  department = 'Development'
```

## Views

- Provide a mechanism to hide certain data from the view of certain users.
- Any relation that is not part of the conceptual model but is visible to a user as a “*virtual relation*” is called a **view**.
- Not physically stored.

## Levels of Abstraction


- Multiple views
- A single Conceptual (Logic) Schema
- A single Physical Schema



## Levels of Abstraction

### Physical Level

- Lowest level
- How the data is physically stored
- It includes
  - Where the data is located
  - File structures
  - Access methods
  - Indexes
- Managed by the **Database Administrator**




## Levels of Abstraction

### Conceptual or Logical Level

- Middle level
- What data is in the DB
- It consists of the schemas described with CREATE TABLE statements
  - Has all the data in the DB


- Has no information on what a user views at external level
- Managed by **Database Designers**



## Levels of Abstraction

### External or View Level

- Highest Level
- Combination of base tables and views
- Views define how certain Users/Groups see data:
  - Full or partial data based on the business requirement
  - Users have different views, based on their levels of access rights
- Exposed to **Users/Applications** and **Database Designers**.




## Data Independence

- A database model exhibits data independence if:


Application programs are protected from changes in the conceptual and physical schemas.

- Why is this important?
  - *Everything changes.*
- Each higher level of the data architecture is immune to changes of the next lower level of the architecture.



## Data Independence Types

- **Physical data independence**
  - Can modify the *physical schema* without causing application programs to be rewritten.
- **Logical data independence**
  - Can modify the *logical schema* without causing application program to be rewritten.



## View Creation and Destruction

- A view is defined using the create view statement which has the form

```
CREATE VIEW view_name AS
< QUERY >
[WITH CHECK OPTION]
```

where < query > is any legal SQL expression.

- Once a view is defined, the view name can be used to refer to the **virtual relation** that the view generates.

# View Creation and Destruction

- **CHECK OPTION**

- Ensures that all UPDATE and INSERTs operations satisfy the condition(s) in the view definition.
- Otherwise, the UPDATE or INSERT returns an error.
- Not implemented in SQLite

- **Destruction**

```
DROP VIEW <view_name>
```

## Views for Security

**Example:**

```
Student(studID, name, address, major, gpa)
```

- This is a view of the Student table without the gpa field.

```
CREATE VIEW SecStudent AS  
SELECT studID, name, address, major  
FROM student
```

## Views for Extensibility

**Example:**

- A company's database includes a relation:  
Part (PartID, weight,...)
- Weight is stored in pounds
- The Company is purchased by a firm that uses metric weights
- Databases must be integrated and use Kg.
  - But old applications use pounds.

# Views for Extensibility

- Solution:
  - Base table with kilograms becomes *MetricPart* for the integrated company

```
CREATE VIEW MetricPart AS
SELECT PartID, 2.2046*weight,      -- no
other changes
FROM Part
```
  - Old programs still call the table *Part*

## Data Partitioning

- Sometimes the data of a database is partitioned.
- **Horizontal:** projection on certain attributes
  - Break up our table based on rows
  - Useful when some attributes are bulky or rarely used
  - Distributed databases
- **Vertical:** selection on certain values (ClientsParis, ClientsLyon)
  - Splitting out extra columns into their own table(s)

## Another Example

- Consider the following relations

Person(name, city)

Purchase(buyer, seller, product, store)

Product(name, maker, category)

```
CREATE VIEW SeattleView AS
```

```
SELECT buyer, seller, product, store
```

```
FROM   Person, Purchase
```

```
WHERE  Person.city = 'Seattle'
```


```
      AND Person.name = Purchase.buyer
```

- We have a new 'virtual table':

SeattleView(buyer, seller, product, store)

## Using (Querying) a View

- Transparency for the user
  - Handled as tables in the database
- Simplify the user's queries
- Useful in architectures client-server



## Using (Querying) a View



## Query using SeattleView

```
SeattleView(buyer, seller, product, store)
Product(name, maker, category)
```

```
SELECT name, store
FROM   SeattleView, Product
WHERE  SeattleView.product = Product.name
      AND Product.category = 'shoes'
```

- When you enter a query that mentions a **view** in the FROM clause, the DBMS expands/rewrites your query to include the view definition.

## View Expansion

- Query using a **view**

```
SELECT name, SeattleView.store
FROM   SeattleView, Product
WHERE  SeattleView.product = Product.name
      AND Product.category = 'shoes'
```

- Expanded query

```
SELECT name, Purchase.store
FROM   Person, Purchase, Product
WHERE  Person.city = 'Seattle'
      AND Person.name = Purchase.buyer
      AND Purchase.product = Product.name
      AND Product.category = 'shoes'
```

## Another Example

- Query using a **view**

```
SELECT buyer, seller
FROM   SeattleView
WHERE  product= 'gizmo'
```

- Expanded query

```
SELECT buyer, seller
FROM   Person, Purchase
WHERE  Person.city = 'Seattle'
       AND Person.name = Purchase.buyer
       AND Purchase.product= 'gizmo'
```

## Views Defined Using Other Views

- One view may be used in the expression defining another view
- A view relation  $v_1$  is said to **depend directly on** a view relation  $v_2$  if  $v_2$  is used in the expression defining  $v_1$
- A view relation  $v_1$  is said to **depend on** view relation  $v_2$  if either  $v_1$  depends directly on  $v_2$  or there is a path of dependencies from  $v_1$  to  $v_2$
- A view relation  $v$  is said to be **recursive** if it depends on itself.

## Updating Views

- How can we insert a tuple into a table that *doesn't exist*?

```
Employee(ssn, name, department, project, salary)

CREATE VIEW Developers AS
  SELECT name, project
  FROM   Employee
  WHERE  department = 'Development'
```

## Updating Views

- How can we insert a tuple into a table that *"doesn't exist"*?

**Example:**

```
Employee(ssn, name, department, project, salary)
Developers(name, project)
```

The following insertion:

```
INSERT INTO Developers
VALUES('Joe', 'Optimizer')
```

becomes:

```
INSERT INTO Employee
VALUES(NULL, 'Joe', NULL, 'Optimizer', NULL)
```

## Non-Updateable Views

- Consider the relations

```
Person(name, city)
Purchase(buyer, seller, product, store)
```

and the view

```
CREATE VIEW SeattleView AS
SELECT seller, product, store
FROM   Person, Purchase
WHERE  Person.city = 'Seattle'
      AND Person.name = Purchase.buyer
```

- How can we add the following tuple to the view?

```
('Joe', 'Shoe Model 12345', 'Nine West')
```

- We need to add 'Joe' to *Person* first.
  - How?!
  - One time?
  - Multiple times?

## Updating Views

- Most SQL implementations allow updates **only on simple views**.
  - The FROM clause has only one database relation.
  - The SELECT clause contains only attribute names of the relation.
    - No expressions, aggregates, or distinct specification.
  - Any attribute not listed in the SELECT clause can be set to NULL.
  - The query does not have a GROUP BY or HAVING clause.

## Updating Views

- SQLite views are read-only and thus you may not be able to execute a DELETE, INSERT or UPDATE statement on a view.
  - A workaround exists
  - Not in the scope of this class
  - This is why WITH CHECK OPTION is not implemented

## Types of Views

- **Virtual views**
  - Used in databases
  - Computed only **on-demand** – *slow at runtime*
  - Always up to date

## Types of Views

- **Materialized views**
  - A physical table containing all the tuples in the result of the query defining the view
  - Used in Data Warehouses (but recently also in DBMS)
  - Precomputed offline – *fast at runtime*

- If relations used in the query are updated, the materialized view result becomes out of date
  - Need to maintain the view, by updating the view whenever the underlying relations are updated.

## Data Warehouse

- A relational database designed for query and analysis rather than for transaction processing.
- Usually contains historical data derived from transaction data.
- Separates analysis workload from transaction workload.
- Enables an organization to consolidate data from several sources.

## Advantages/Disadvantages of Views

ADVANTAGES	DISADVANTAGES
Data independence	Update restriction
Currency	Structure restriction
Improved security	Performance
Reduced complexity	
Convenience	
Customization	
Data integrity	

## Summary

- A view is a stored query definition
- Views can be very useful
  - Privacy

- Easier query writing
- Extensibility
- Not all views can be updated unambiguously
- Three levels of abstraction in a relational DBMS
  - Yields data independence, logical and physical

## Access Rights in SQL

- The SQL security model
- Granting and revoking privileges

## Discretionary Access Control

- Each user is given appropriate access rights (*privileges*) on specific DB objects
- Explicit grant of rights on objects to individuals.
- Users obtain certain privileges when they create an object
  - Can *pass* some or all of these privileges to other users **at their discretion**
- Although flexible, can be circumvented by devious unauthorized user tricking an authorized user into revealing sensitive data.

## Terminology

- **Privacy** Users should not be able to see and use data they are not supposed to.

e.g., A student can't see other students' grades.

- **Security** No one should be able to enter the system and / or impact its behavior without being authorized to do so.

e.g., Delete or change data without being authorized

- **Integrity** Authorized users should not be able to modify things they are not supposed to.

e.g., Only instructors can assign grades.

- **Availability** Users should be able to see and modify things they are allowed to.

e.g. The DB should always be operational

## SQL Security Related Terminology

- **User**
  - Not the schema object, just a name for a session of an individual user
  - Identification by Authorization ID (user name)
- **Role**
  - Name for a role, to which rights may be assigned
  - May be granted to users / applications
- **Privileges (Rights)**
  - System privileges
  - Object (data) privileges: creator has all privileges
- **Operations**
  - GRANT < privilege >
  - REVOKE < privilege >

## Roles and Users

- **Roles** define a set of privileges for a (potentially) large set of **Users**

```
CREATE ROLE sales_people;  
-- grant some privileges to sales_people  
-- grant sales_people role to users
```

- Much more economic than direct privileges
- Roles may be assigned to roles
- Often assigned to applications instead of individual users

## Privileges

- Right to perform SQL statement type on objects
- Assigned to users or roles (authorization IDs)
- *Creator of object*: all privileges for that object
- *Administrator*: management of system privileges

## Privileges

- The privileges defined by the ISO standard:
  - SELECT - retrieve data from a table
  - INSERT - insert new rows into a table
  - UPDATE - modify rows of data in a table
  - DELETE - delete rows of data from a table
  - REFERENCE - reference columns of a named table in integrity constraints
  - USAGE - use domains, collations, character sets, and translations

## Grant Privileges

- Syntax

```
GRANT <privileges> ON <object>  
TO [<users>|<role>]  
[WITH GRANT OPTION]
```

- GRANT OPTION: Right to pass privilege on to other users
  - Only owner can execute CREATE, ALTER, and DROP



**Example:** Privilege to INSERT particular columns in a table

```
GRANT INSERT
ON <tablename(<attributenames>)>
TO <users>
[WITH GRANT OPTION]
```

- Access matrix : < user > has < right > on < object >

## Examples

```
GRANT INSERT, SELECT ON Movie TO Klaus
```

- Klaus can query **'Movie'** or insert tuples into it.

```
GRANT DELETE ON Movie TO shop_owner WITH GRANT OPTION
```

- Anna can delete **'Movie'** tuples, and also authorize others to do so

```
GRANT UPDATE (price_Day) ON Movie TO movie_staff
```

- Staff can update (only) the price field of **'Movie'** tuples

```
GRANT SELECT ON MovieView TO Customers
```

- This does NOT allow the customers to query **'Movie'** directly!

# Revoke Privileges

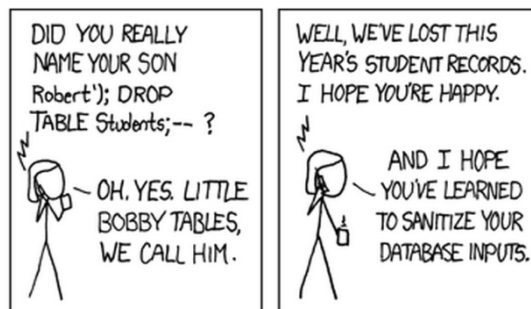
- Syntax

```
REVOKE <privileges>  
ON <object>  
FROM <users>  
[RESTRICT | CASCADE]
```

- RESTRICT: only revoke if none of the privileges have been granted by these users.
- CASCADE: revoke from all users that have been granted the privilege by these users.
- Privilege given from different users must be revoked from all users to loose privilege.

## Summary

- Security of DB and their applications is extremely important.
- Roles make privileges with many users manageable.
- Views also play an important role.
- Fine granular access restriction on objects is very important.



<https://www.xkcd.com/327/>