

Reshaping and pivot tables


Reshaping by pivoting DataFrame objects

Pivot

`df.pivot(index='foo', columns='bar', values='baz')`

df

	foo	bar	baz	zoo
0	one	A	1	x
1	one	B	2	y
2	one	C	3	z
3	two	A	4	q
4	two	B	5	w
5	two	C	6	t



bar	A	B	C
foo			
one	1	2	3
two	4	5	6

Data is often stored in so-called “stacked” or “record” format:

```
In [1]: df
Out[1]:
```

	date	variable	value
0	2000-01-03	A	0.469112
1	2000-01-04	A	-0.282863
2	2000-01-05	A	-1.509059
3	2000-01-03	B	-1.135632
4	2000-01-04	B	1.212112
5	2000-01-05	B	-0.173215
6	2000-01-03	C	0.119209
7	2000-01-04	C	-1.044236
8	2000-01-05	C	-0.861849
9	2000-01-03	D	-2.104569
10	2000-01-04	D	-0.494929
11	2000-01-05	D	1.071804

For the curious here is how the above DataFrame was created:

```
import pandas.util.testing as tm

tm.N = 3

def unpivot(frame):
    N, K = frame.shape
    data = {'value': frame.to_numpy().ravel('F'),
            'variable': np.asarray(frame.columns).repeat(N),
            'date': np.tile(np.asarray(frame.index), K)}
    return pd.DataFrame(data, columns=['date', 'variable', 'value'])
```

```
df = unpivot(tm.makeTimeDataFrame())
```

To select out everything for variable A we could do:

```
In [2]: df[df['variable'] == 'A']
Out[2]:
```

	date	variable	value
0	2000-01-03	A	0.469112
1	2000-01-04	A	-0.282863
2	2000-01-05	A	-1.509059

But suppose we wish to do time series operations with the variables. A better representation would be where the columns are the unique variables and an index of dates identifies individual observations. To reshape the data into this form, we use the `DataFrame.pivot()` method (also implemented as a top level function `pivot()`):

```
In [3]: df.pivot(index='date', columns='variable', values='value')
Out[3]:
```

variable	A	B	C	D
date				
2000-01-03	0.469112	-1.135632	0.119209	-2.104569
2000-01-04	-0.282863	1.212112	-1.044236	-0.494929
2000-01-05	-1.509059	-0.173215	-0.861849	1.071804

If the `values` argument is omitted, and the input `DataFrame` has more than one column of values which are not used as column or index inputs to `pivot`, then the resulting “pivoted” `DataFrame` will have [hierarchical columns](#) whose topmost level indicates the respective value column:

```
In [4]: df['value2'] = df['value'] * 2
In [5]: pivoted = df.pivot(index='date', columns='variable')
In [6]: pivoted
Out[6]:
```

	value				value2			
variable	A	B	C	D	A	B	C	
date								
2000-01-03	0.469112	-1.135632	0.119209	-2.104569	0.938225	-2.271265	0.238417	-4
2000-01-04	-0.282863	1.212112	-1.044236	-0.494929	-0.565727	2.424224	-2.088472	-0
2000-01-05	-1.509059	-0.173215	-0.861849	1.071804	-3.018117	-0.346429	-1.723698	2

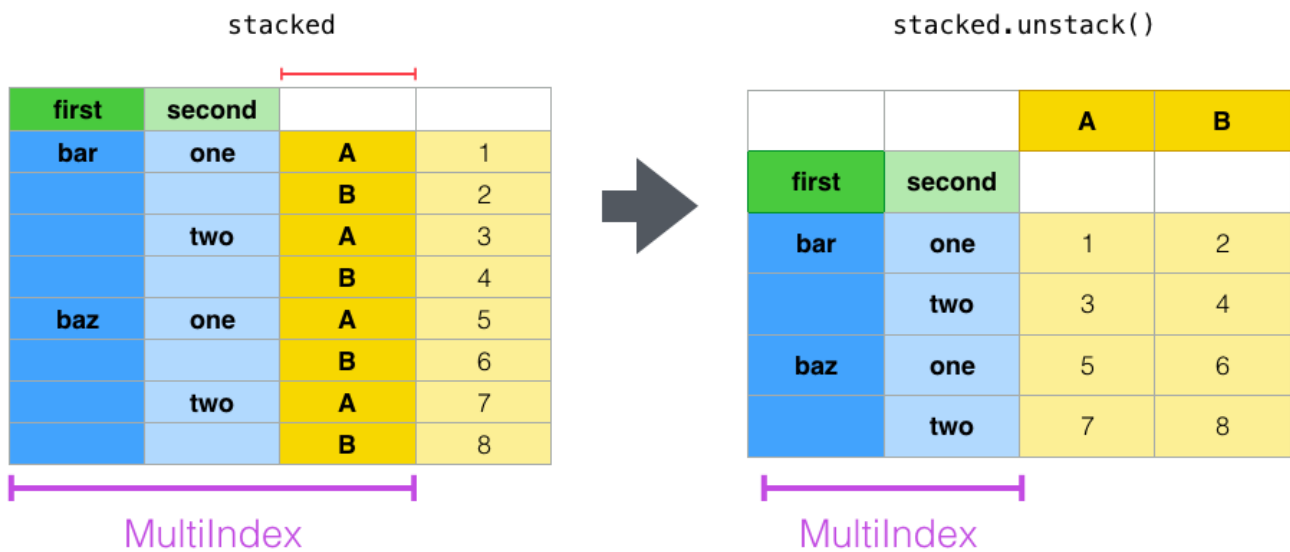
You can then select subsets from the pivoted `DataFrame`:

```
In [7]: pivoted['value2']
Out[7]:
```

variable	A	B	C	D
date				
2000-01-03	0.938225	-2.271265	0.238417	-4.209138
2000-01-04	-0.565727	2.424224	-2.088472	-0.989859
2000-01-05	-3.018117	-0.346429	-1.723698	2.143608

Note that this returns a view on the underlying data in the case where the data are homogeneously-typed.

Unstack



The clearest way to explain is by example. Let's take a prior example data set from the hierarchical indexing section:

```
In [8]: tuples = list(zip(*[['bar', 'bar', 'baz', 'baz',
...:                        'foo', 'foo', 'qux', 'qux'],
...:                      ['one', 'two', 'one', 'two',
...:                      'one', 'two', 'one', 'two']]))

In [9]: index = pd.MultiIndex.from_tuples(tuples, names=['first', 'second'])

In [10]: df = pd.DataFrame(np.random.randn(8, 2), index=index, columns=['A', 'B'])

In [11]: df2 = df[:4]

In [12]: df2
Out[12]:
```

		A	B
first	second		
bar	one	0.721555	-0.706771
	two	-1.039575	0.271860
baz	one	-0.424972	0.567020
	two	0.276232	-1.087401

The stack function “compresses” a level in the DataFrame’s columns to produce either:

- A Series, in the case of a simple column Index.
- A DataFrame, in the case of a MultiIndex in the columns.

If the columns have a MultiIndex, you can choose which level to stack. The stacked level becomes the new lowest level in a MultiIndex on the columns:

```
In [13]: stacked = df2.stack()

In [14]: stacked
Out[14]:
```

```

first second
bar  one    A    0.721555
      one    B   -0.706771
      two    A   -1.039575
      two    B    0.271860
baz   one    A   -0.424972
      one    B    0.567020
      two    A    0.276232
      two    B   -1.087401
dtype: float64

```

With a “stacked” DataFrame or Series (having a MultiIndex as the index), the inverse operation of stack is unstack, which by default unstacks the **last level**:

```

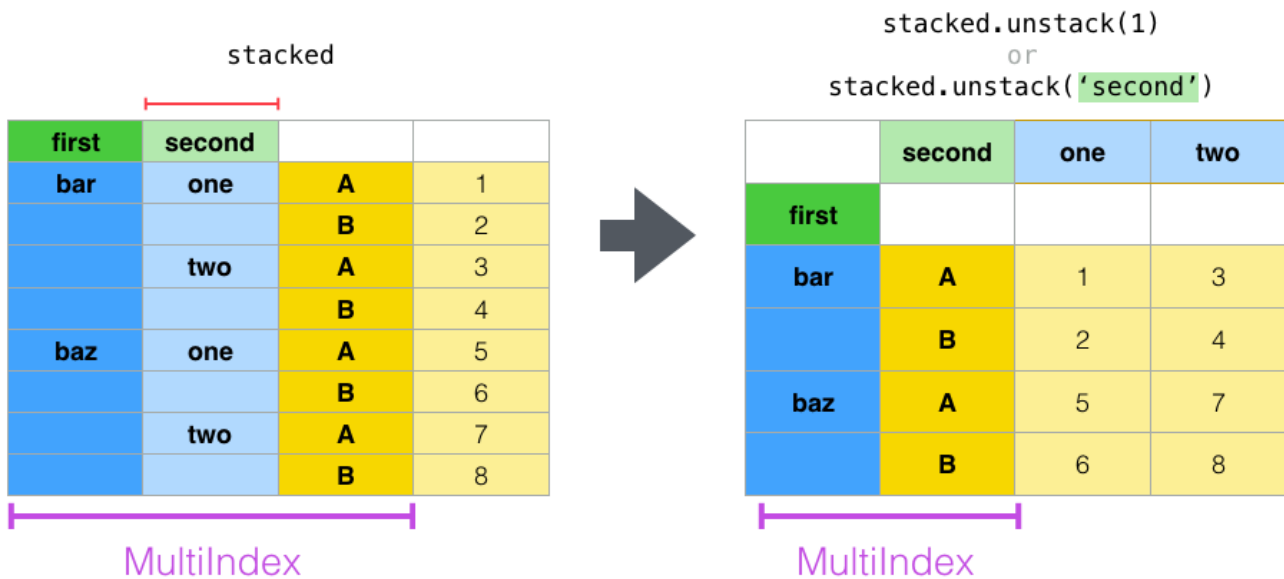
In [15]: stacked.unstack()
Out[15]:
              A          B
first second
bar  one    0.721555 -0.706771
      two   -1.039575  0.271860
baz   one   -0.424972  0.567020
      two    0.276232 -1.087401

In [16]: stacked.unstack(1)
Out[16]:
second      one      two
first
bar  A  0.721555 -1.039575
      B -0.706771  0.271860
baz  A -0.424972  0.276232
      B  0.567020 -1.087401

In [17]: stacked.unstack(0)
Out[17]:
first      bar      baz
second
one  A  0.721555 -0.424972
      B -0.706771  0.567020
two  A -1.039575  0.276232
      B  0.271860 -1.087401

```

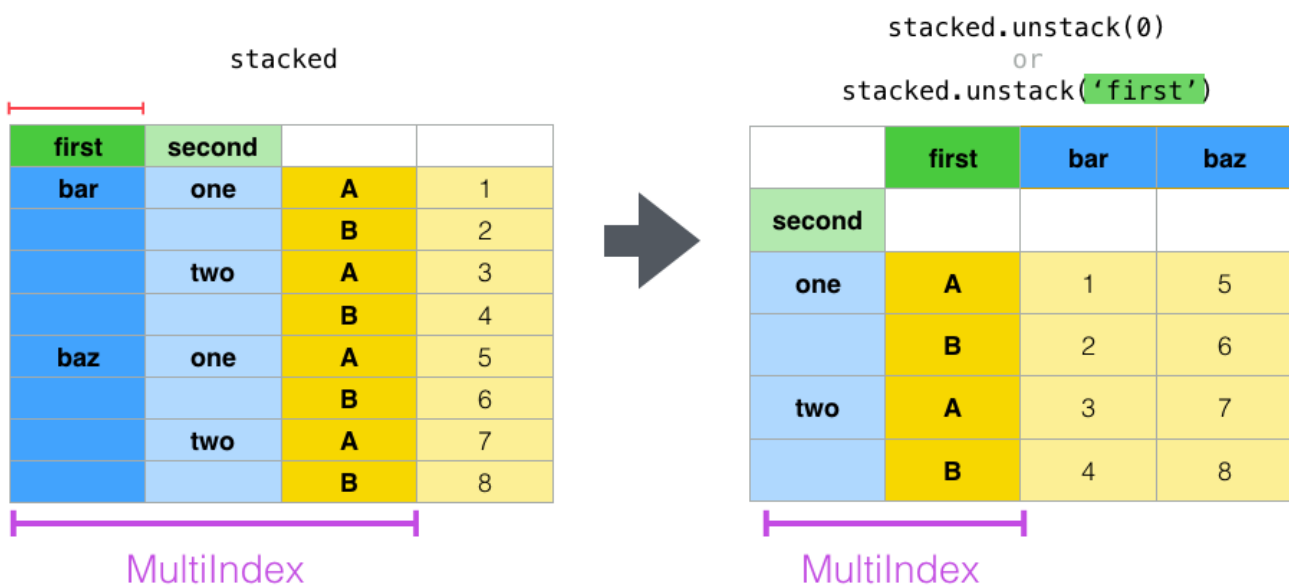
Unstack(1)



If the indexes have names, you can use the level names instead of specifying the level numbers:

```
In [18]: stacked.unstack('second')
Out[18]:
second      one      two
first
bar   A  0.721555 -1.039575
      B -0.706771  0.271860
baz   A -0.424972  0.276232
      B  0.567020 -1.087401
```

Unstack(0)



Notice that the stack and unstack methods implicitly sort the index levels involved. Hence a call to stack and then unstack, or vice versa, will result in a **sorted** copy of the original DataFrame or Series:

```
In [19]: index = pd.MultiIndex.from_product([[2, 1], ['a', 'b']])
In [20]: df = pd.DataFrame(np.random.randn(4), index=index, columns=['A'])
In [21]: df
Out[21]:
           A
2 a -0.370647
  b -1.157892
1 a -1.344312
  b  0.844885

In [22]: all(df.unstack().stack() == df.sort_index())
Out[22]: True
```

The above code will raise a `TypeError` if the call to `sort_index` is removed.

Multiple levels

You may also stack or unstack more than one level at a time by passing a list of levels, in which case the end result is as if each level in the list were processed individually.

```
In [23]: columns = pd.MultiIndex.from_tuples([
.....:     ('A', 'cat', 'long'), ('B', 'cat', 'long'),
.....:     ('A', 'dog', 'short'), ('B', 'dog', 'short')],
.....:     names=['exp', 'animal', 'hair_length'])
.....:

In [24]: df = pd.DataFrame(np.random.randn(4, 4), columns=columns)

In [25]: df
Out[25]:
exp          A          B          A          B
animal      cat      cat      dog      dog
hair_length long    long    short    short
0          1.075770 -0.109050  1.643563 -1.469388
1          0.357021 -0.674600 -1.776904 -0.968914
2         -1.294524  0.413738  0.276662 -0.472035
3         -0.013960 -0.362543 -0.006154 -0.923061

In [26]: df.stack(level=['animal', 'hair_length'])
Out[26]:
exp          A          B
animal hair_length
0 cat      long      1.075770 -0.109050
  dog      short      1.643563 -1.469388
1 cat      long      0.357021 -0.674600
  dog      short     -1.776904 -0.968914
2 cat      long     -1.294524  0.413738
  dog      short      0.276662 -0.472035
3 cat      long     -0.013960 -0.362543
  dog      short     -0.006154 -0.923061
```

The list of levels can contain either level names or level numbers (but not a mixture of the two).

```
# df.stack(level=['animal', 'hair_length'])
# from above is equivalent to:
In [27]: df.stack(level=[1, 2])
Out[27]:
exp          A          B
```

```

    animal hair_length
0 cat      long      1.075770 -0.109050
  dog      short      1.643563 -1.469388
1 cat      long      0.357021 -0.674600
  dog      short     -1.776904 -0.968914
2 cat      long     -1.294524  0.413738
  dog      short      0.276662 -0.472035
3 cat      long     -0.013960 -0.362543
  dog      short     -0.006154 -0.923061

```

Missing data

These functions are intelligent about handling missing data and do not expect each subgroup within the hierarchical index to have the same set of labels. They also can handle the index being unsorted (but you can make it sorted by calling `sort_index`, of course). Here is a more complex example:

```

In [28]: columns = pd.MultiIndex.from_tuples([('A', 'cat'), ('B', 'dog'),
.....:                                     ('B', 'cat'), ('A', 'dog')],
.....:                                     names=['exp', 'animal'])

In [29]: index = pd.MultiIndex.from_product([('bar', 'baz', 'foo', 'qux'),
.....:                                     ('one', 'two')],
.....:                                     names=['first', 'second'])

In [30]: df = pd.DataFrame(np.random.randn(8, 4), index=index, columns=columns)

In [31]: df2 = df.iloc[[0, 1, 2, 4, 5, 7]]

In [32]: df2
Out[32]:
exp      A      B      A
animal   cat   dog   cat   dog
first second
bar  one    0.895717  0.805244 -1.206412  2.565646
     two    1.431256  1.340309 -1.170299 -0.226169
baz   one    0.410835  0.813850  0.132003 -0.827317
foo   one   -1.413681  1.607920  1.024180  0.569605
     two    0.875906 -2.211372  0.974466 -2.006747
qux   two   -1.226825  0.769804 -1.281247 -0.727707

```

As mentioned above, `stack` can be called with a `level` argument to select which level in the columns to stack:

```

In [33]: df2.stack('exp')
Out[33]:
animal      cat      dog
first second exp
bar  one    A    0.895717  2.565646
     two    B   -1.206412  0.805244
baz   one    A    0.410835 -0.827317
     two    B    0.132003  0.813850
foo   one    A   -1.413681  0.569605
     two    B    1.024180  1.607920
qux   two    A    0.875906 -2.006747
     two    B    0.974466 -2.211372

```



```
In [34]: df2.stack('animal')
Out[34]:
```

exp			A	B
bar	one	cat	0.895717	-1.206412
		dog	2.565646	0.805244
	two	cat	1.431256	-1.170299
		dog	-0.226169	1.340309
baz	one	cat	0.410835	0.132003
		dog	-0.827317	0.813850
foo	one	cat	-1.413681	1.024180
		dog	0.569605	1.607920
	two	cat	0.875906	0.974466
		dog	-2.006747	-2.211372
qux	two	cat	-1.226825	-1.281247
		dog	-0.727707	0.769804

Unstacking can result in missing values if subgroups do not have the same set of labels. By default, missing values will be replaced with the default fill value for that data type, NaN for float, NaT for datetimelike, etc. For integer types, by default data will be converted to float and missing values will be set to NaN.

```
In [35]: df3 = df.iloc[[0, 1, 4, 7], [1, 2]]
In [36]: df3
Out[36]:
```

exp			B	
animal	first	second	dog	cat
bar	one		0.805244	-1.206412
	two		1.340309	-1.170299
foo	one		1.607920	1.024180
qux	two		0.769804	-1.281247

```
In [37]: df3.unstack()
Out[37]:
```

exp			B		
animal	second	first	dog	cat	
			one	two	one
bar			one	two	one
foo			one	two	one
qux			one	two	one

New in version 0.18.0.

Alternatively, unstack takes an optional `fill_value` argument, for specifying the value of missing data.

```
In [38]: df3.unstack(fill_value=-1e9)
Out[38]:
```

exp			B		
animal	second	first	dog	cat	
			one	two	one
bar			one	two	one
foo			one	two	one
qux			one	two	one

With a MultiIndex

Unstacking when the columns are a MultiIndex is also careful about doing the right thing:

```
In [39]: df[:3].unstack(0)
Out[39]:
exp      A      B      A
animal   cat   dog
first    bar   bar   baz   baz   cat   baz   dog
second
one      0.895717 0.410835 0.805244 0.81385 -1.206412 0.132003 2.565646 -0.827317
two      1.431256      NaN 1.340309      NaN -1.170299      NaN -0.226169      NaN

In [40]: df2.unstack(1)
Out[40]:
exp      A      B      A
animal   cat   dog
second   one   two   one   two   cat   two   dog
first
bar      0.895717 1.431256 0.805244 1.340309 -1.206412 -1.170299 2.565646 -0.226169
baz      0.410835      NaN 0.813850      NaN 0.132003      NaN -0.827317
foo     -1.413681 0.875906 1.607920 -2.211372 1.024180 0.974466 0.569605 -2.006705
qux      NaN -1.226825      NaN 0.769804      NaN -1.281247      NaN -0.727705
```

Reshaping by Melt

Melt

df3					df3.melt(id_vars=['first', 'last'])				
	first	last	height	weight		first	last	variable	value
0	John	Doe	5.5	130	0	John	Doe	height	5.5
1	Mary	Bo	6.0	150	1	Mary	Bo	height	6.0
					2	John	Doe	weight	130
					3	Mary	Bo	weight	150

The top-level `melt()` function and the corresponding `DataFrame.melt()` are useful to massage a `DataFrame` into a format where one or more columns are *identifier variables*, while all other columns, considered *measured variables*, are “unpivoted” to the row axis, leaving just two non-identifier columns, “variable” and “value”. The names of those columns can be customized by supplying the `var_name` and `value_name` parameters.

For instance,

```
In [41]: cheese = pd.DataFrame({'first': ['John', 'Mary'],
.....:                          'last': ['Doe', 'Bo'],
.....:                          'height': [5.5, 6.0],
.....:                          'weight': [130, 150]})
```

```

In [42]: cheese
Out[42]:
   first last  height  weight
0  John  Doe     5.5     130
1  Mary   Bo     6.0     150

In [43]: cheese.melt(id_vars=['first', 'last'])
Out[43]:
   first last variable  value
0  John  Doe   height     5.5
1  Mary   Bo   height     6.0
2  John  Doe  weight    130.0
3  Mary   Bo  weight    150.0

In [44]: cheese.melt(id_vars=['first', 'last'], var_name='quantity')
Out[44]:
   first last quantity  value
0  John  Doe   height     5.5
1  Mary   Bo   height     6.0
2  John  Doe  weight    130.0
3  Mary   Bo  weight    150.0

```

Another way to transform is to use the `wide_to_long()` panel data convenience function. It is less flexible than `melt()`, but more user-friendly.

```

In [45]: dft = pd.DataFrame({"A1970": {0: "a", 1: "b", 2: "c"},
.....:                      "A1980": {0: "d", 1: "e", 2: "f"},
.....:                      "B1970": {0: 2.5, 1: 1.2, 2: .7},
.....:                      "B1980": {0: 3.2, 1: 1.3, 2: .1},
.....:                      "X": dict(zip(range(3), np.random.randn(3)))
.....:                      })

In [46]: dft["id"] = dft.index

In [47]: dft
Out[47]:
   A1970 A1980  B1970  B1980         X  id
0      a      d     2.5     3.2 -0.121306  0
1      b      e     1.2     1.3 -0.097883  1
2      c      f     0.7     0.1  0.695775  2

In [48]: pd.wide_to_long(dft, ["A", "B"], i="id", j="year")
Out[48]:
         X  A    B
id year
0  1970 -0.121306  a  2.5
1  1970 -0.097883  b  1.2
2  1970  0.695775  c  0.7
0  1980 -0.121306  d  3.2
1  1980 -0.097883  e  1.3
2  1980  0.695775  f  0.1

```

Combining with stats and GroupBy

It should be no shock that combining `pivot / stack / unstack` with `GroupBy` and the basic `Series` and `DataFrame` statistical functions can produce some very expressive and fast data manipulations.

```

In [49]: df
Out[49]:

```

```

exp          A          B
animal       cat       dog
first second
bar  one    0.895717  0.805244 -1.206412  2.565646
     two    1.431256  1.340309 -1.170299 -0.226169
baz   one    0.410835  0.813850  0.132003 -0.827317
     two   -0.076467 -1.187678  1.130127 -1.436737
foo   one   -1.413681  1.607920  1.024180  0.569605
     two    0.875906 -2.211372  0.974466 -2.006747
qux   one   -0.410001 -0.078638  0.545952 -1.219217
     two   -1.226825  0.769804 -1.281247 -0.727707

```

```
In [50]: df.stack().mean(1).unstack()
```

```
Out[50]:
```

```

animal       cat       dog
first second
bar  one   -0.155347  1.685445
     two    0.130479  0.557070
baz   one    0.271419 -0.006733
     two    0.526830 -1.312207
foo   one   -0.194750  1.088763
     two    0.925186 -2.109060
qux   one    0.067976 -0.648927
     two   -1.254036  0.021048

```

```
# same result, another way
```

```
In [51]: df.groupby(level=1, axis=1).mean()
```

```
Out[51]:
```

```

animal       cat       dog
first second
bar  one   -0.155347  1.685445
     two    0.130479  0.557070
baz   one    0.271419 -0.006733
     two    0.526830 -1.312207
foo   one   -0.194750  1.088763
     two    0.925186 -2.109060
qux   one    0.067976 -0.648927
     two   -1.254036  0.021048

```

```
In [52]: df.stack().groupby(level=1).mean()
```

```
Out[52]:
```

```

exp          A          B
second
one    0.071448  0.455513
two   -0.424186 -0.204486

```

```
In [53]: df.mean().unstack(0)
```

```
Out[53]:
```

```

exp          A          B
animal
cat    0.060843  0.018596
dog   -0.413580  0.232430

```

Pivot tables

While `pivot()` provides general purpose pivoting with various data types (strings, numerics, etc.), pandas also provides `pivot_table()` for pivoting with aggregation of numeric data.

The function `pivot_table()` can be used to create spreadsheet-style pivot tables. See the [cookbook](#) for some advanced strategies.

It takes a number of arguments:

- `data`: a `DataFrame` object.

- `values`: a column or a list of columns to aggregate.
- `index`: a column, Grouper, array which has the same length as data, or list of them. Keys to group by on the pivot table index. If an array is passed, it is being used as the same manner as column values.
- `columns`: a column, Grouper, array which has the same length as data, or list of them. Keys to group by on the pivot table column. If an array is passed, it is being used as the same manner as column values.
- `aggfunc`: function to use for aggregation, defaulting to `numpy.mean`.

Consider a data set like this:

```
In [54]: import datetime

In [55]: df = pd.DataFrame({'A': ['one', 'one', 'two', 'three'] * 6,
.....:                    'B': ['A', 'B', 'C'] * 8,
.....:                    'C': ['foo', 'foo', 'foo', 'bar', 'bar', 'bar'] * 4,
.....:                    'D': np.random.randn(24),
.....:                    'E': np.random.randn(24),
.....:                    'F': [datetime.datetime(2013, i, 1) for i in range(1, 13)]
.....:                    + [datetime.datetime(2013, i, 15) for i in range(1, 13)]})

In [56]: df
Out[56]:
```

	A	B	C	D	E	F
0	one	A	foo	0.341734	-0.317441	2013-01-01
1	one	B	foo	0.959726	-1.236269	2013-02-01
2	two	C	foo	-1.110336	0.896171	2013-03-01
3	three	A	bar	-0.619976	-0.487602	2013-04-01
4	one	B	bar	0.149748	-0.082240	2013-05-01
...
19	three	B	foo	0.690579	-2.213588	2013-08-15
20	one	C	foo	0.995761	1.063327	2013-09-15
21	one	A	bar	2.396780	1.266143	2013-10-15
22	two	B	bar	0.014871	0.299368	2013-11-15
23	three	C	bar	3.357427	-0.863838	2013-12-15

[24 rows x 6 columns]

We can produce pivot tables from this data very easily:

```
In [57]: pd.pivot_table(df, values='D', index=['A', 'B'], columns=['C'])
Out[57]:
```

		bar	foo
A	B		
one	A	1.120915	-0.514058
	B	-0.338421	0.002759
	C	-0.538846	0.699535
three	A	-1.181568	NaN
	B	NaN	0.433512
	C	0.588783	NaN
two	A	NaN	1.000985
	B	0.158248	NaN
	C	NaN	0.176180

```
In [58]: pd.pivot_table(df, values='D', index=['B'], columns=['A', 'C'], aggfunc=np
Out[58]:
```

		one	three	two
		bar	bar	bar
B	A	foo	foo	foo
A	one	2.241830	-1.028115	-2.363137
	three	NaN	NaN	NaN
	two	NaN	NaN	2.001971

```

D -0.070843  0.005518  NaN  0.007024  0.310493  NaN
C -1.077692  1.399070  1.177566  NaN  NaN  0.352360

In [59]: pd.pivot_table(df, values=['D', 'E'], index=['B'], columns=['A', 'C'],
.....:                  aggfunc=np.sum)
.....:
Out[59]:

```

		D				E		
		one	foo	three		one	foo	two
		bar		bar		bar		bar
	B							
A	2.241830	-1.028115	-2.363137	NaN	NaN	2.001971	2.786113	-0.043211
B	-0.676843	0.005518	NaN	0.867024	0.316495	NaN	1.368280	-1.103384
C	-1.077692	1.399070	1.177566	NaN	NaN	0.352360	-1.976883	1.495717

The result object is a DataFrame having potentially hierarchical indexes on the rows and columns. If the values column name is not given, the pivot table will include all of the data that can be aggregated in an additional level of hierarchy in the columns:

```

In [60]: pd.pivot_table(df, index=['A', 'B'], columns=['C'])
Out[60]:

```

		D		E	
		bar	foo	bar	foo
	B				
one	A	1.120915	-0.514058	1.393057	-0.021605
	B	-0.338421	0.002759	0.684140	-0.551692
	C	-0.538846	0.699535	-0.988442	0.747859
three	A	-1.181568	NaN	0.961289	NaN
	B	NaN	0.433512	NaN	-1.064372
	C	0.588783	NaN	-0.131830	NaN
two	A	NaN	1.000985	NaN	0.064245
	B	0.158248	NaN	-0.097147	NaN
	C	NaN	0.176180	NaN	0.436241

Also, you can use Grouper for index and columns keywords. For detail of Grouper, see [Grouping with a Grouper specification](#).

```

In [61]: pd.pivot_table(df, values='D', index=pd.Grouper(freq='M', key='F'),
.....:                  columns='C')
.....:
Out[61]:

```

		bar	foo
	F		
2013-01-31	NaN	-0.514058	
2013-02-28	NaN	0.002759	
2013-03-31	NaN	0.176180	
2013-04-30	-1.181568	NaN	
2013-05-31	-0.338421	NaN	
2013-06-30	-0.538846	NaN	
2013-07-31	NaN	1.000985	
2013-08-31	NaN	0.433512	
2013-09-30	NaN	0.699535	
2013-10-31	1.120915	NaN	
2013-11-30	0.158248	NaN	
2013-12-31	0.588783	NaN	

You can render a nice output of the table omitting the missing values by calling `to_string` if you wish:

```
In [62]: table = pd.pivot_table(df, index=['A', 'B'], columns=['C'])
```

```
In [63]: print(table.to_string(na_rep=''))
```

		D		E	
C		bar	foo	bar	foo
one	A	1.120915	-0.514058	1.393057	-0.021605
	B	-0.338421	0.002759	0.684140	-0.551692
	C	-0.538846	0.699535	-0.988442	0.747859
three	A	-1.181568		0.961289	
	B		0.433512		-1.064372
	C	0.588783		-0.131830	
two	A		1.000985		0.064245
	B	0.158248		-0.097147	
	C		0.176180		0.436241

Note that `pivot_table` is also available as an instance method on `DataFrame`,
i.e. `DataFrame.pivot_table()`.

Adding margins

If you pass `margins=True` to `pivot_table`, special `All` columns and rows will be added with partial group aggregates across the categories on the rows and columns:

```
In [64]: df.pivot_table(index=['A', 'B'], columns='C', margins=True, aggfunc=np.std)
Out[64]:
```

		D		E	
C		bar	foo	bar	foo
one	A	1.804346	1.210272	1.569879	0.179483
	B	0.690376	1.353355	0.898998	1.083825
	C	0.273641	0.418926	0.771139	1.689271
three	A	0.794212	NaN	0.794212	2.049040
	B	NaN	0.363548	0.363548	NaN
	C	3.915454	NaN	3.915454	1.035215
two	A	NaN	0.442998	0.442998	NaN
	B	0.202765	NaN	0.202765	0.560757
	C	NaN	1.819408	1.819408	NaN
All		1.556686	0.952552	1.246608	1.250924

Cross tabulations

Use `crosstab()` to compute a cross-tabulation of two (or more) factors. By default `crosstab` computes a frequency table of the factors unless an array of values and an aggregation function are passed.

It takes a number of arguments

- `index`: array-like, values to group by in the rows.
- `columns`: array-like, values to group by in the columns.
- `values`: array-like, optional, array of values to aggregate according to the factors.
- `aggfunc`: function, optional, If no values array is passed, computes a frequency table.
- `rownames`: sequence, default `None`, must match number of row arrays passed.
- `colnames`: sequence, default `None`, if passed, must match number of column arrays passed.
- `margins`: boolean, default `False`, Add row/column margins (subtotals)

- `normalize`: boolean, {'all', 'index', 'columns'}, or {0,1}, default False. Normalize by dividing all values by the sum of values.

Any Series passed will have their name attributes used unless row or column names for the cross-tabulation are specified

For example:

```
In [65]: foo, bar, dull, shiny, one, two = 'foo', 'bar', 'dull', 'shiny', 'one', 'two'
In [66]: a = np.array([foo, foo, bar, bar, foo, foo], dtype=object)
In [67]: b = np.array([one, one, two, one, two, one], dtype=object)
In [68]: c = np.array([dull, dull, shiny, dull, dull, shiny], dtype=object)
In [69]: pd.crosstab(a, [b, c], rownames=['a'], colnames=['b', 'c'])
Out[69]:
b      one      two
c  dull shiny dull shiny
a
bar    1     0     0     1
foo    2     1     1     0
```

If `crosstab` receives only two Series, it will provide a frequency table.

```
In [70]: df = pd.DataFrame({'A': [1, 2, 2, 2, 2], 'B': [3, 3, 4, 4, 4],
.....:                    'C': [1, 1, np.nan, 1, 1]})
.....:

In [71]: df
Out[71]:
   A  B    C
0  1  3  1.0
1  2  3  1.0
2  2  4  NaN
3  2  4  1.0
4  2  4  1.0

In [72]: pd.crosstab(df.A, df.B)
Out[72]:
B    3    4
A
1    1    0
2    1    3
```

Any input passed containing `Categorical` data will have **all** of its categories included in the cross-tabulation, even if the actual data does not contain any instances of a particular category.

```
In [73]: foo = pd.Categorical(['a', 'b'], categories=['a', 'b', 'c'])
In [74]: bar = pd.Categorical(['d', 'e'], categories=['d', 'e', 'f'])
In [75]: pd.crosstab(foo, bar)
Out[75]:
col_0  d  e
row_0
a      1  0
b      0  1
```


Normalization

New in version 0.18.1.

Frequency tables can also be normalized to show percentages rather than counts using the `normalize` argument:

```
In [76]: pd.crosstab(df.A, df.B, normalize=True)
Out[76]:
B      3      4
A
1    0.2    0.0
2    0.2    0.6
```

`normalize` can also normalize values within each row or within each column:

```
In [77]: pd.crosstab(df.A, df.B, normalize='columns')
Out[77]:
B      3      4
A
1    0.5    0.0
2    0.5    1.0
```

`crosstab` can also be passed a third Series and an aggregation function (`aggfunc`) that will be applied to the values of the third Series within each group defined by the first two Series:

```
In [78]: pd.crosstab(df.A, df.B, values=df.C, aggfunc=np.sum)
Out[78]:
B      3      4
A
1    1.0    NaN
2    1.0    2.0
```

Adding margins

Finally, one can also add margins or normalize this output.

```
In [79]: pd.crosstab(df.A, df.B, values=df.C, aggfunc=np.sum, normalize=True,
.....:               margins=True)
Out[79]:
B      3      4    All
A
1    0.25  0.0  0.25
2    0.25  0.5  0.75
All  0.50  0.5  1.00
```

Tiling

The `cut()` function computes groupings for the values of the input array and is often used to transform continuous variables to discrete or categorical variables:

```
In [80]: ages = np.array([10, 15, 13, 12, 23, 25, 28, 59, 60])

In [81]: pd.cut(ages, bins=3)
Out[81]:
[(9.95, 26.667], (9.95, 26.667], (9.95, 26.667], (9.95, 26.667], (9.95, 26.667], (9
Categories (3, interval[float64]): [(9.95, 26.667] < (26.667, 43.333] < (43.333, 60]
```

If the bins keyword is an integer, then equal-width bins are formed. Alternatively we can specify custom bin-edges:

```
In [82]: c = pd.cut(ages, bins=[0, 18, 35, 70])

In [83]: c
Out[83]:
[(0, 18], (0, 18], (0, 18], (0, 18], (18, 35], (18, 35], (18, 35], (35, 70], (35, 70]
Categories (3, interval[int64]): [(0, 18] < (18, 35] < (35, 70]]
```

New in version 0.20.0.

If the bins keyword is an IntervalIndex, then these will be used to bin the passed data.:

```
pd.cut([25, 20, 50], bins=c.categories)
```

Computing indicator / dummy variables

To convert a categorical variable into a “dummy” or “indicator” DataFrame, for example a column in a DataFrame (a Series) which has k distinct values, can derive a DataFrame containing k columns of 1s and 0s using `get_dummies()`:

```
In [84]: df = pd.DataFrame({'key': list('bbacab'), 'data1': range(6)})

In [85]: pd.get_dummies(df['key'])
Out[85]:
   a  b  c
0  0  1  0
1  0  1  0
2  1  0  0
3  0  0  1
4  1  0  0
5  0  1  0
```

Sometimes it’s useful to prefix the column names, for example when merging the result with the original DataFrame:

```
In [86]: dummies = pd.get_dummies(df['key'], prefix='key')

In [87]: dummies
Out[87]:
   key_a  key_b  key_c
0      0      1      0
1      0      1      0
2      1      0      0
```

```

3      0      0      1
4      1      0      0
5      0      1      0

In [88]: df[['data1']].join(dummies)
Out[88]:
   data1  key_a  key_b  key_c
0      0      0      1      0
1      1      0      1      0
2      2      1      0      0
3      3      0      0      1
4      4      1      0      0
5      5      0      1      0

```

This function is often used along with discretization functions like `cut`:

```

In [89]: values = np.random.randn(10)

In [90]: values
Out[90]:
array([ 0.4082, -1.0481, -0.0257, -0.9884,  0.0941,  1.2627,  1.29   ,
        0.0824, -0.0558,  0.5366])

In [91]: bins = [0, 0.2, 0.4, 0.6, 0.8, 1]

In [92]: pd.get_dummies(pd.cut(values, bins))
Out[92]:
   (0.0, 0.2]  (0.2, 0.4]  (0.4, 0.6]  (0.6, 0.8]  (0.8, 1.0]
0             0           0           1           0           0
1             0           0           0           0           0
2             0           0           0           0           0
3             0           0           0           0           0
4             1           0           0           0           0
5             0           0           0           0           0
6             0           0           0           0           0
7             1           0           0           0           0
8             0           0           0           0           0
9             0           0           1           0           0

```

See also [Series.str.get_dummies](#).

`get_dummies()` also accepts a `DataFrame`. By default all categorical variables (categorical in the statistical sense, those with *object* or *categorical* dtype) are encoded as dummy variables.

```

In [93]: df = pd.DataFrame({'A': ['a', 'b', 'a'], 'B': ['c', 'c', 'b'],
.....:                    'C': [1, 2, 3]})
.....:

In [94]: pd.get_dummies(df)
Out[94]:
   C  A_a  A_b  B_b  B_c
0  1    1    0    0    1
1  2    0    1    0    1
2  3    1    0    1    0

```

All non-object columns are included untouched in the output. You can control the columns that are encoded with the `columns` keyword.

```

In [95]: pd.get_dummies(df, columns=['A'])
Out[95]:

```

	B	C	A_a	A_b
0	c	1	1	0
1	c	2	0	1
2	b	3	1	0

Notice that the B column is still included in the output, it just hasn't been encoded. You can drop B before calling `get_dummies` if you don't want to include it in the output.

As with the Series version, you can pass values for the `prefix` and `prefix_sep`. By default the column name is used as the prefix, and `'_'` as the prefix separator. You can specify `prefix` and `prefix_sep` in 3 ways:

- string: Use the same value for `prefix` or `prefix_sep` for each column to be encoded.
- list: Must be the same length as the number of columns being encoded.
- dict: Mapping column name to prefix.

```
In [96]: simple = pd.get_dummies(df, prefix='new_prefix')
```

```
In [97]: simple
```

```
Out[97]:
```

	C	new_prefix_a	new_prefix_b	new_prefix_b	new_prefix_c
0	1	1	0	0	1
1	2	0	1	0	1
2	3	1	0	1	0

```
In [98]: from_list = pd.get_dummies(df, prefix=['from_A', 'from_B'])
```

```
In [99]: from_list
```

```
Out[99]:
```

	C	from_A_a	from_A_b	from_B_b	from_B_c
0	1	1	0	0	1
1	2	0	1	0	1
2	3	1	0	1	0

```
In [100]: from_dict = pd.get_dummies(df, prefix={'B': 'from_B', 'A': 'from_A'})
```

```
In [101]: from_dict
```

```
Out[101]:
```

	C	from_A_a	from_A_b	from_B_b	from_B_c
0	1	1	0	0	1
1	2	0	1	0	1
2	3	1	0	1	0

New in version 0.18.0.

Sometimes it will be useful to only keep $k-1$ levels of a categorical variable to avoid collinearity when feeding the result to statistical models. You can switch to this mode by turn on `drop_first`.

```
In [102]: s = pd.Series(list('abcaa'))
```

```
In [103]: pd.get_dummies(s)
```

```
Out[103]:
```

	a	b	c
0	1	0	0
1	0	1	0
2	0	0	1
3	1	0	0
4	1	0	0

```
In [104]: pd.get_dummies(s, drop_first=True)
```

```
Out[104]:
   b  c
0  0  0
1  1  0
2  0  1
3  0  0
4  0  0
```

When a column contains only one level, it will be omitted in the result.

```
In [105]: df = pd.DataFrame({'A': list('aaaaa'), 'B': list('ababc')})

In [106]: pd.get_dummies(df)
Out[106]:
   A_a  B_a  B_b  B_c
0     1    1    0    0
1     1    0    1    0
2     1    1    0    0
3     1    0    1    0
4     1    0    0    1

In [107]: pd.get_dummies(df, drop_first=True)
Out[107]:
   B_b  B_c
0     0    0
1     1    0
2     0    0
3     1    0
4     0    1
```

By default new columns will have `np.uint8` dtype. To choose another dtype, use the `dtype` argument:

```
In [108]: df = pd.DataFrame({'A': list('abc'), 'B': [1.1, 2.2, 3.3]})

In [109]: pd.get_dummies(df, dtype=bool).dtypes
Out[109]:
B          float64
A_a         bool
A_b         bool
A_c         bool
dtype: object
```

New in version 0.23.0.

Factorizing values

To encode 1-d values as an enumerated type use `factorize()`:

```
In [110]: x = pd.Series(['A', 'A', np.nan, 'B', 3.14, np.inf])

In [111]: x
Out[111]:
0      A
1      A
2     NaN
3      B
4    3.14
5     inf
dtype: object
```

```
In [112]: labels, uniques = pd.factorize(x)

In [113]: labels
Out[113]: array([ 0,  0, -1,  1,  2,  3])

In [114]: uniques
Out[114]: Index(['A', 'B', 3.14, inf], dtype='object')
```

Note that `factorize` is similar to `numpy.unique`, but differs in its handling of NaN:

Note: The following `numpy.unique` will fail under Python 3 with a `TypeError` because of an ordering bug. See also [here](#).

```
In [1]: x = pd.Series(['A', 'A', np.nan, 'B', 3.14, np.inf])
In [2]: pd.factorize(x, sort=True)
Out[2]: (array([ 2,  2, -1,  3,  0,  1]),
         Index([3.14, inf, 'A', 'B'], dtype='object'))

In [3]: np.unique(x, return_inverse=True)[::-1]
Out[3]: (array([3, 3, 0, 4, 1, 2]), array([nan, 3.14, inf, 'A', 'B'], dtype=object))
```

Note: If you just want to handle one column as a categorical variable (like R's `factor`), you can use `df["cat_col"] = pd.Categorical(df["col"])` or `df["cat_col"] = df["col"].astype("category")`. For full docs on [Categorical](#), see the [Categorical introduction](#) and the [API documentation](#).

Examples

In this section, we will review frequently asked questions and examples. The column names and relevant column values are named to correspond with how this `DataFrame` will be pivoted in the answers below.

```
In [115]: np.random.seed([3, 1415])

In [116]: n = 20

In [117]: cols = np.array(['key', 'row', 'item', 'col'])

In [118]: df = cols + pd.DataFrame((np.random.randint(5, size=(n, 4))
.....:                               // [2, 1, 2, 1]).astype(str))
.....:

In [119]: df.columns = cols

In [120]: df = df.join(pd.DataFrame(np.random.rand(n, 2).round(2)).add_prefix('val'))

In [121]: df
Out[121]:
   key  row  item  col  val0  val1
0  key0 row3  item1 col3  0.81  0.04
1  key1 row2  item1 col2  0.44  0.07
2  key1 row0  item1 col0  0.77  0.01
3  key0 row4  item0 col2  0.15  0.59
4  key1 row0  item2 col1  0.81  0.64
..  ...  ...  ...  ...  ...  ...
```

```

15  key0  row3  item1  col1  0.31  0.23
16  key0  row0  item2  col3  0.86  0.01
17  key0  row4  item0  col3  0.64  0.21
18  key2  row2  item2  col0  0.13  0.45
19  key0  row2  item0  col4  0.37  0.70

[20 rows x 6 columns]

```

Pivoting with single aggregations

Suppose we wanted to pivot `df` such that the `col` values are columns, `row` values are the index, and the mean of `val0` are the values? In particular, the resulting DataFrame should look like:

Note: col col0 col1 col2 col3 col4 row row0 0.77 0.605 NaN 0.860 0.65 row2 0.13 NaN 0.395 0.500 0.25 row3 NaN 0.310 NaN 0.545 NaN row4 NaN 0.100 0.395 0.760 0.24

This solution uses `pivot_table()`. Also note that `aggfunc='mean'` is the default. It is included here to be explicit.

```

In [122]: df.pivot_table(
.....:     values='val0', index='row', columns='col', aggfunc='mean')
.....:
Out[122]:
col  col0  col1  col2  col3  col4
row
row0  0.77  0.605   NaN  0.860  0.65
row2  0.13   NaN  0.395  0.500  0.25
row3   NaN  0.310   NaN  0.545   NaN
row4   NaN  0.100  0.395  0.760  0.24

```

Note that we can also replace the missing values by using the `fill_value` parameter.

```

In [123]: df.pivot_table(
.....:     values='val0', index='row', columns='col', aggfunc='mean', fill_value=0)
.....:
Out[123]:
col  col0  col1  col2  col3  col4
row
row0  0.77  0.605  0.000  0.860  0.65
row2  0.13  0.000  0.395  0.500  0.25
row3  0.00  0.310  0.000  0.545  0.00
row4  0.00  0.100  0.395  0.760  0.24

```

Also note that we can pass in other aggregation functions as well. For example, we can also pass in `sum`.

```

In [124]: df.pivot_table(
.....:     values='val0', index='row', columns='col', aggfunc='sum', fill_value=0)
.....:
Out[124]:
col  col0  col1  col2  col3  col4
row
row0  0.77  1.21  0.00  0.86  0.65
row2  0.13  0.00  0.79  0.50  0.50

```

```
row3  0.00  0.31  0.00  1.09  0.00
row4  0.00  0.10  0.79  1.52  0.24
```

Another aggregation we can do is calculate the frequency in which the columns and rows occur together a.k.a. “cross tabulation”. To do this, we can pass size to the aggfunc parameter.

```
In [125]: df.pivot_table(index='row', columns='col', fill_value=0, aggfunc='size')
Out[125]:
```

col	col0	col1	col2	col3	col4
row					
row0	1	2	0	1	1
row2	1	0	2	1	2
row3	0	1	0	2	0
row4	0	1	2	2	1

Pivoting with multiple aggregations

We can also perform multiple aggregations. For example, to perform both a sum and mean, we can pass in a list to the aggfunc argument.

```
In [126]: df.pivot_table(
.....:     values='val0', index='row', columns='col', aggfunc=['mean', 'sum'])
Out[126]:
```

	mean					sum				
col	col0	col1	col2	col3	col4	col0	col1	col2	col3	col4
row										
row0	0.77	0.605	NaN	0.860	0.65	0.77	1.21	NaN	0.86	0.65
row2	0.13	NaN	0.395	0.500	0.25	0.13	NaN	0.79	0.50	0.50
row3	NaN	0.310	NaN	0.545	NaN	NaN	0.31	NaN	1.09	NaN
row4	NaN	0.100	0.395	0.760	0.24	NaN	0.10	0.79	1.52	0.24

Note to aggregate over multiple value columns, we can pass in a list to the values parameter.

```
In [127]: df.pivot_table(
.....:     values=['val0', 'val1'], index='row', columns='col', aggfunc=['mean'])
Out[127]:
```

	mean					val1				
col	col0	col1	col2	col3	col4	col0	col1	col2	col3	col4
row										
row0	0.77	0.605	NaN	0.860	0.65	0.01	0.745	NaN	0.010	0.02
row2	0.13	NaN	0.395	0.500	0.25	0.45	NaN	0.34	0.440	0.79
row3	NaN	0.310	NaN	0.545	NaN	NaN	0.230	NaN	0.075	NaN
row4	NaN	0.100	0.395	0.760	0.24	NaN	0.070	0.42	0.300	0.46

Note to subdivide over multiple columns we can pass in a list to the columns parameter.

```
In [128]: df.pivot_table(
.....:     values=['val0'], index='row', columns=['item', 'col'], aggfunc=['mean'])
Out[128]:
```

	mean
val0	

	item0			item1					item2			
col	col2	col3	col4	col0	col1	col2	col3	col4	col0	col1	col3	col4
row												
row0	NaN	NaN	NaN	0.77	NaN	NaN	NaN	NaN	NaN	0.605	0.86	0.65
row2	0.35	NaN	0.37	NaN	NaN	0.44	NaN	NaN	0.13	NaN	0.50	0.13
row3	NaN	NaN	NaN	NaN	0.31	NaN	0.81	NaN	NaN	NaN	0.28	NaN
row4	0.15	0.64	NaN	NaN	0.10	0.64	0.88	0.24	NaN	NaN	NaN	NaN

Exploding a list-like column

New in version 0.25.0.

Sometimes the values in a column are list-like.

```
In [129]: keys = ['panda1', 'panda2', 'panda3']
In [130]: values = [['eats', 'shoots'], ['shoots', 'leaves'], ['eats', 'leaves']]
In [131]: df = pd.DataFrame({'keys': keys, 'values': values})
In [132]: df
Out[132]:
```

	keys	values
0	panda1	[eats, shoots]
1	panda2	[shoots, leaves]
2	panda3	[eats, leaves]

We can ‘explode’ the values column, transforming each list-like to a separate row, by using `explode()`. This will replicate the index values from the original row:

```
In [133]: df['values'].explode()
Out[133]:
```

0	eats
0	shoots
1	shoots
1	leaves
2	eats
2	leaves

Name: values, dtype: object

You can also explode the column in the DataFrame.

```
In [134]: df.explode('values')
Out[134]:
```

	keys	values
0	panda1	eats
0	panda1	shoots
1	panda2	shoots
1	panda2	leaves
2	panda3	eats
2	panda3	leaves

`Series.explode()` will replace empty lists with `np.nan` and preserve scalar entries. The dtype of the resulting Series is always object.

```

In [135]: s = pd.Series([[1, 2, 3], 'foo', [], ['a', 'b']])

In [136]: s
Out[136]:
0    [1, 2, 3]
1         foo
2          []
3    [a, b]
dtype: object

In [137]: s.explode()
Out[137]:
0      1
0      2
0      3
1    foo
2   NaN
3     a
3     b
dtype: object

```

Here is a typical usecase. You have comma separated strings in a column and want to expand this.

```

In [138]: df = pd.DataFrame([{'var1': 'a,b,c', 'var2': 1},
.....:                      {'var1': 'd,e,f', 'var2': 2}])
.....:

In [139]: df
Out[139]:
   var1  var2
0  a,b,c     1
1  d,e,f     2

```

Creating a long form DataFrame is now straightforward using explode and chained operations

```

In [140]: df.assign(var1=df.var1.str.split(',')).explode('var1')
Out[140]:
   var1  var2
0     a     1
0     b     1
0     c     1
1     d     2
1     e     2
1     f     2

```