

Deep Learning

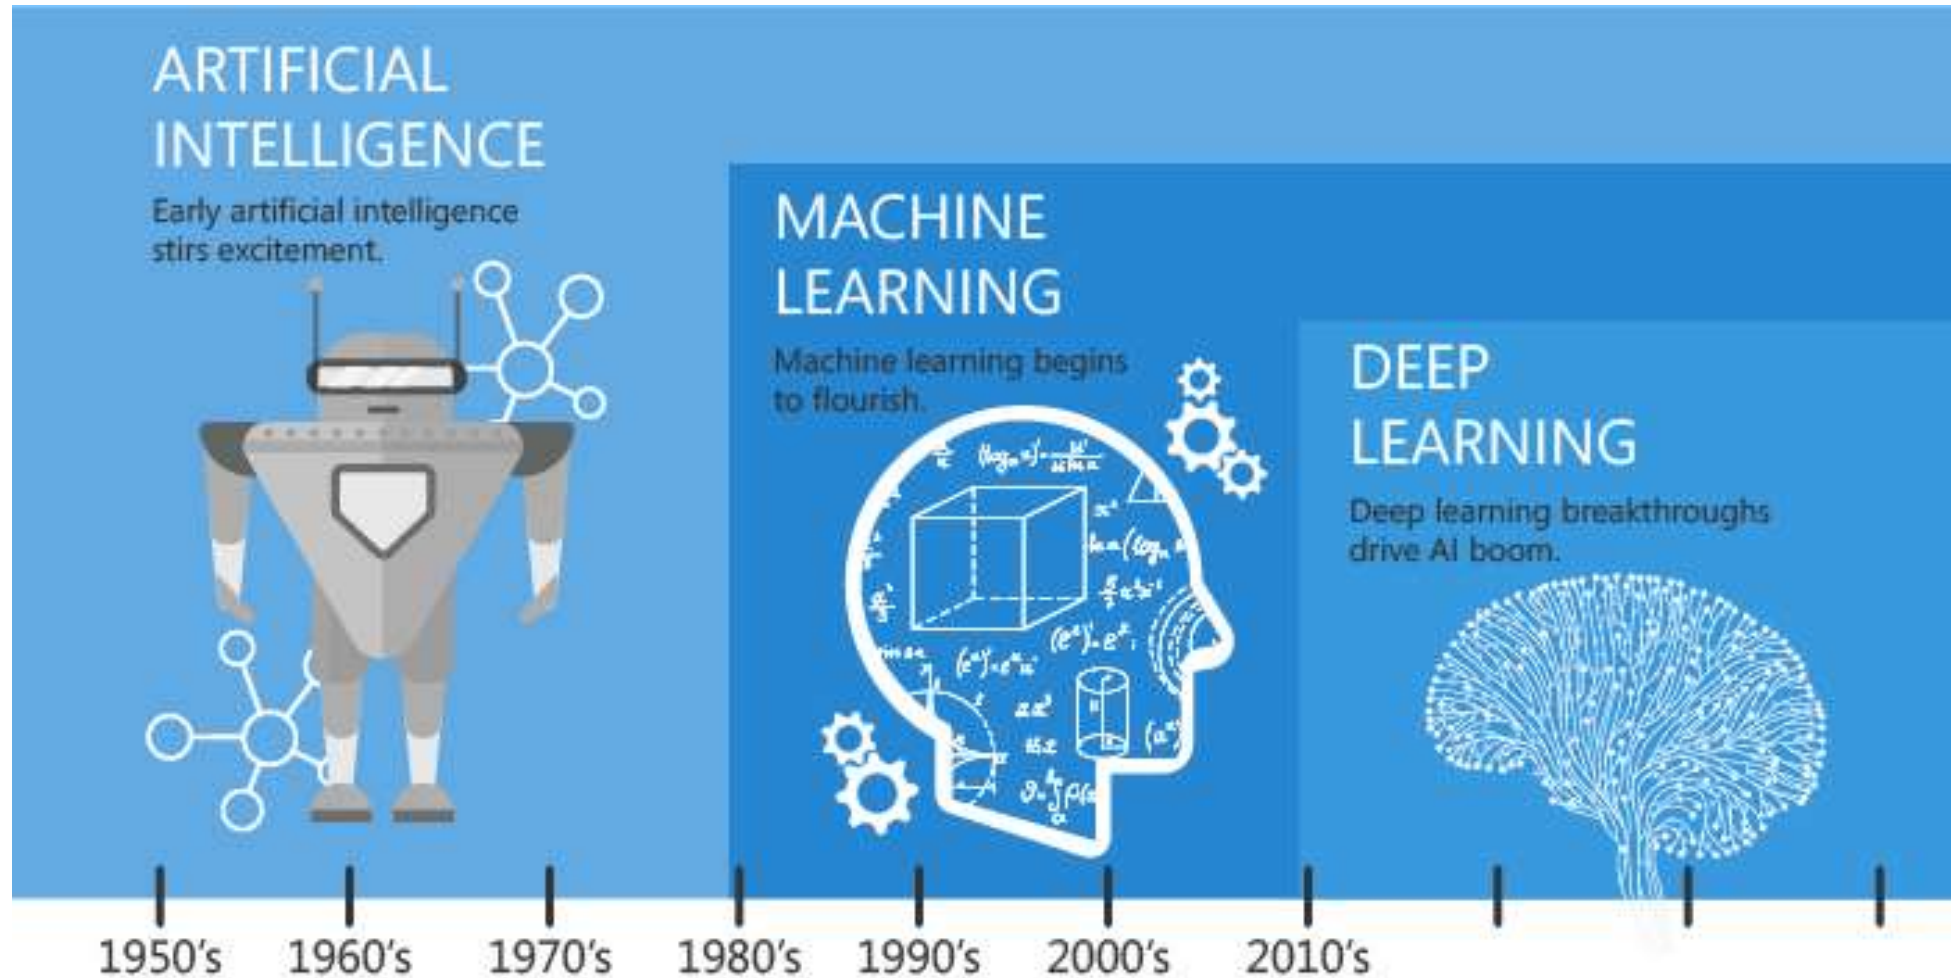
Stéphane Gentric, Ph. D.
Senior Expert
Research Unit Manager



Fev 2020

- **Intro**
- **Deep @ Idemia**
- **Data**
- **Learning Frameworks**
- **Neural Network**
- **CNN Architecture**
- **Optimization**
- **Overfitting**
- **More DNN Architectures**
- **Adversaries**
- **Tips**

ARTIFICIAL INTELLIGENCE, MACHINE LEARNING, AND DEEP LEARNING



Since an early flush of optimism in the 1950's, smaller subsets of artificial intelligence - first machine learning, then deep learning, a subset of machine learning - have created ever larger disruptions.




In “Nature” 27 January 2016:

“DeepMind’s program AlphaGo beat Fan Hui, the European Go champion, five times out of five in tournament conditions...”

“AlphaGo was not preprogrammed to play Go: rather, it learned using a general-purpose algorithm that allowed it to interpret the game’s patterns.”

“...AlphaGo program applied **deep learning** in neural networks (convolutional NN) — brain-inspired programs in which connections between layers of simulated neurons are strengthened through examples and experience.”

ALPHAGO ZERO: LEARNING FROM SCRATCH





The ImageNet challenge


- **Crucial in demonstrating the effectiveness of deep CNNs**
- **Problem: recognize object categories in Internet imagery**
- **The 2012 ImageNet Large Scale Visual Recognition Challenge (ILSVRC) classification task - classify image from Flickr and other search engines into 1 of 1000 possible object categories**
- **Serves as a standard benchmark for deep learning**
- **The imagery was hand-labeled based on the presence or absence of an object belonging to these categories**
- **There are 1.2 million images in the training set with 732-1300 training images available per class**
- **A random subset of 50,000 images was used as the validation set, and 100,000 images were used for the test set where there are 50 and 100 images per class respectively**



IMAGENET CHALLENGE: RECOGNIZE OBJECT CATEGORIES



- 1000 object categories
- **Hand-labeled** images
- 1.2 million images in the training set





A plateau, then rapid advances

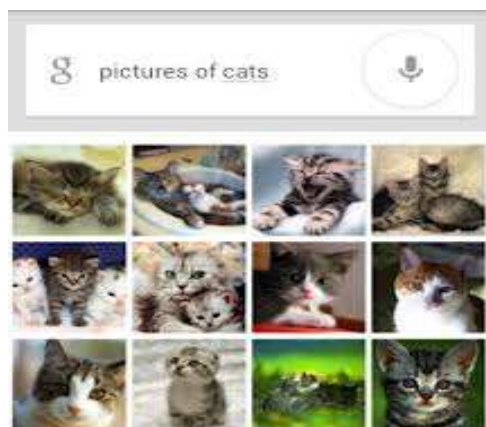
“Top-5 error” is the % of times that the target label does not appear among the 5 highest-probability predictions

Visual recognition methods not based on deep CNNs hit a plateau in performance at 25%

Name	Layers	Top-5 Error (%)	References
AlexNet	8	15.3	Krizhevsky et al. (2012)
VGG Net	19	7.3	Simonyan and Zisserman (2014)
ResNet	152	3.6	He et al. (2016)



Deep Learning – from Research to Technology



Deep Learning - breakthrough in visual
and speech recognition



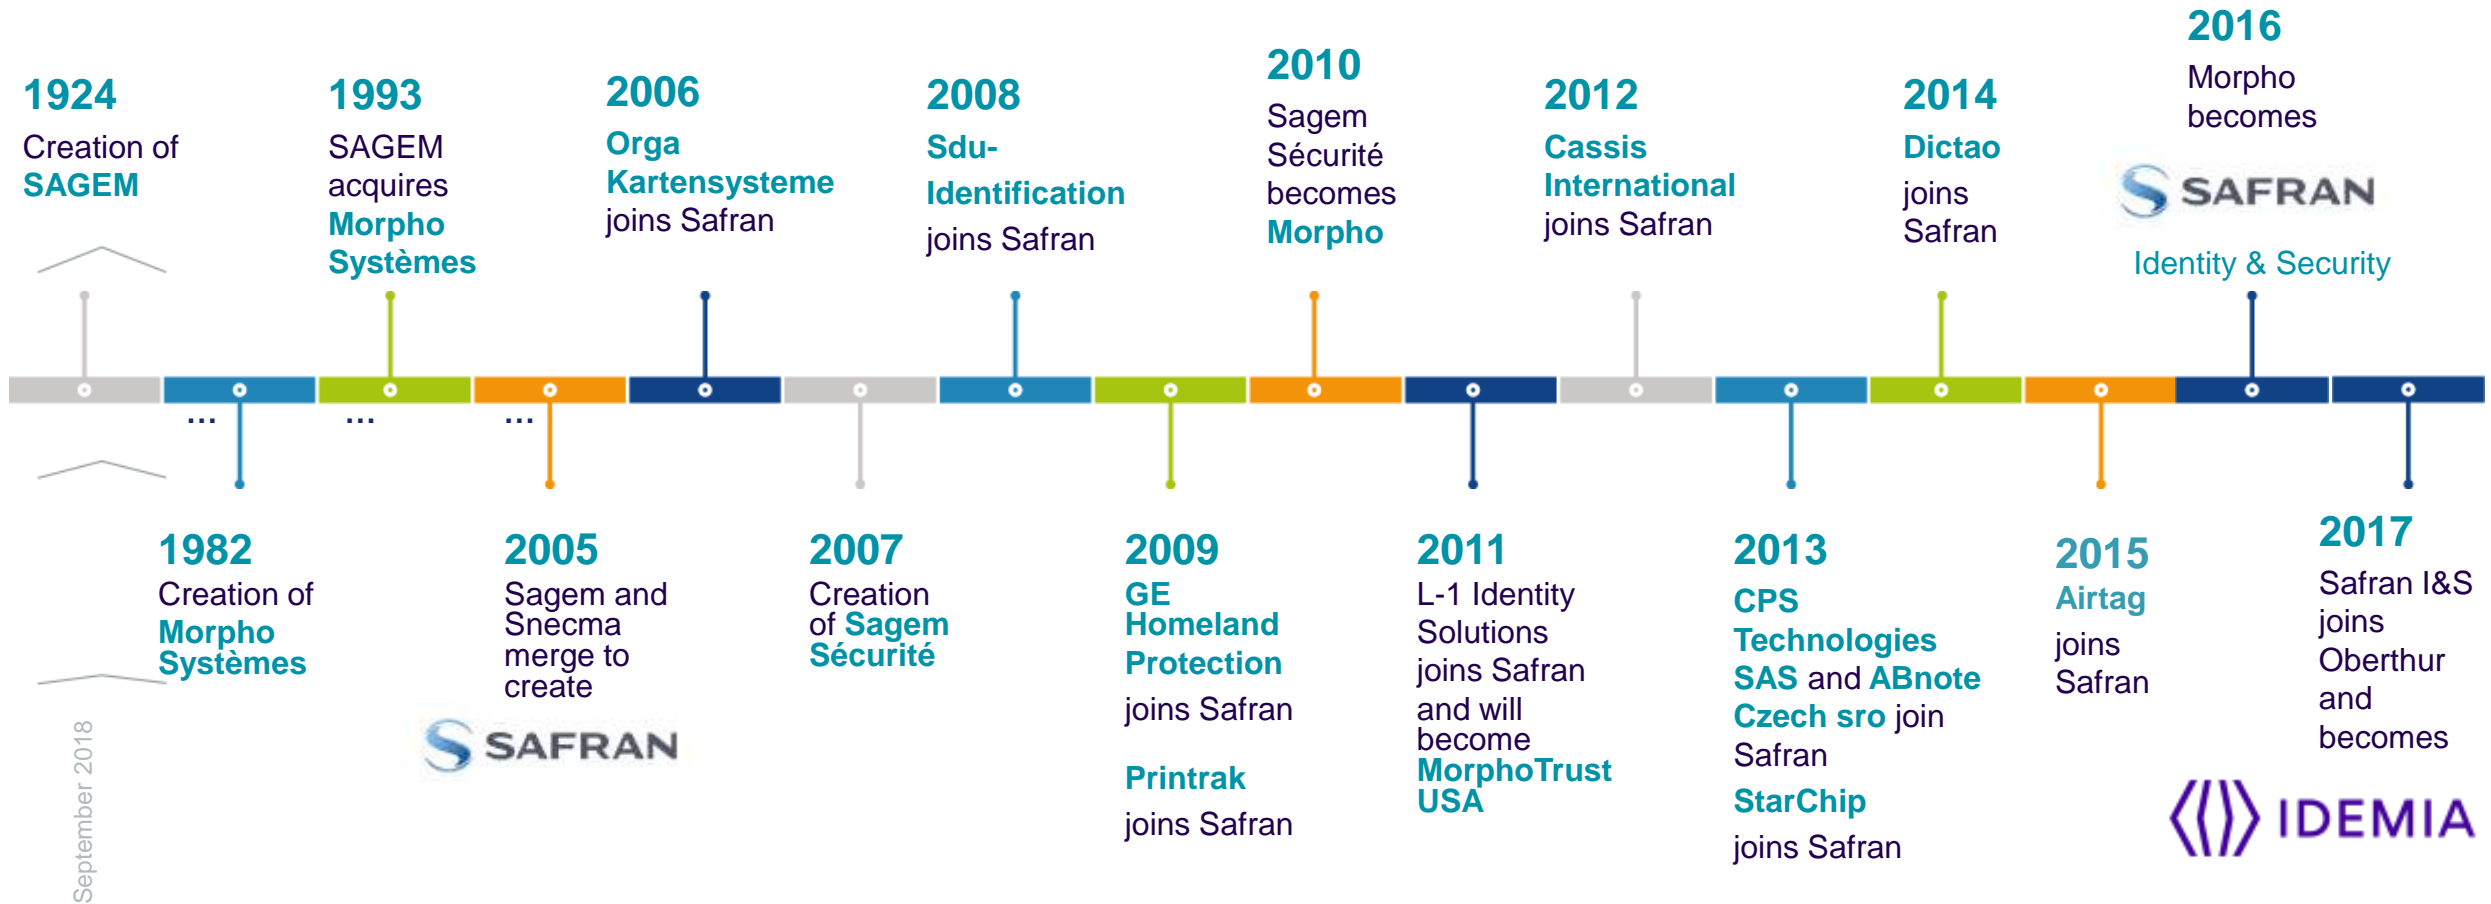
Deep @ Idemia



2



A short history of Biometrics at Idemia





#1 Leader de l'Identité Augmentée dans un monde de plus en plus digital



Près de
3 Milliards de \$

de chiffre d'affaires



+ de 14 000 Collaborateurs

dont + de 2 000 en R&D



Plus de

80 Nationalités



**Un ensemble de solutions
innovantes et intégrées**

200 Millions d'\$

en R&D en 2017 & plus de 1400
familles de brevets actives



**Production de cartes
à grande échelle**



+ de 3 Milliards

de documents d'identité délivrés
dans le monde



1,2 Milliard

de cartes SIM livrées

en 2017



+ de 850 Millions

de cartes de paiement
produites en 2017



+ de 250 Millions

d'éléments sécurisés embarqués déployés
dans le monde



+ de 20

solutions de gestion d'abonnement
déployées dans le monde



N°1 des systèmes
d'identification biométriques

1 800 institutions
financières nous font
confiance

N°1 des solutions
d'identité civile

+ de 500


opérateurs mobiles nous
font confiance

**Les plus grands OEM
industriels**

nous font confiance

N°1 dans l'émission de permis
de conduire aux États-Unis

Global R&D - Research and Technology Unit



4 fingers



Compact finger





Finger OTF



3D Face



Iris
At distance





Face Recognition Markets

Global R&D / R&T Unit Overview

- **Identity Solution**
 - Secure Deliverance of a unique right (Vote, ID document)
 - Statistically good Quality images ... but very large population
- **Forensic Application**
 - Heterogeneous image quality
 - Police officer may spend time on a single case
- **Border Control Application**
 - Increase workflow & Maintain Security level
 - Passenger cooperation level.
- **Connected Objects & Financial Institutions**
 - Liveness, UI, low CPU/Memory

UIDAI

- Unique ID for India
- 10 fingers + 2 irises + 1 face
- 1+ billion citizens

FBI

- central US system for latent fingerprints, tenprints and faces

SmartGate

- Australian/New Zealand Airports
- 80+ Gates - facial
- 40+ millions crossing

March 2019



Object & Face Recognition Pipeline

Static Image

Find the Objects/Faces in images. Mugshot, ID Doc, Selfie, in the wild

Understand and correct : pose, light sources, occlusions

Extract discriminative features and qualities

Optimize speed / accuracy trade-off

Detection
Tracking

Registration

Features
Extraction

Matching

Video Stream

Track objects/persons in video
Limit false face detections.

Track and refine pose estimation to improve robustness.

Extract discriminative features and qualities
Select and fuse features.

Control False Alarm Rate.




Face Detection Algorithm

Classifier


- › Robust features, ex: Haar, Filter bank, SIFT, SURF, HOG, color histograms
- › Machine learning, ex: SVM, Ada-Float boost, decision tree

The classifiers are learned on a database of faces and non faces. Bootstrap and selection of hard examples.
The detection pipeline is a fusion of those approaches for obtaining the best speed/accuracy tradeoff.






Traditional Approaches



These methods are composed of two steps



1. Agnostic features are extracted from the images

- › Most of the discriminative information is born by the gradients
- › Various representations of the gradients have been proposed
 - SIFT, HOG, LBP
 - Filter bank responses (Gaussian derivatives, Gabor, etc.)



2. A transformation of the initial feature space is learned

- › It reduces the dimensionality and concentrates the discriminative information (remove noise and redundancy)
- › Most methods simply use a linear transformation
 - LDA, ITML, LMNN, ...
- › But more complex methods exist (e.g. local metric)




A well-engineered traditional method gives accurate results on well aligned images but lacks of robustness



Deep Networks for Face Recognition

- A Face template is learned directly from the pixels.



2-3 week on GPUs
millions of iterations
millions of images

Biometric Performances comes from :


- The size of the image database (thousands of identities, hundreds of images per id)
- The architecture of the Network (number of layers, maps ...)
- The loss functions



Deep Networks for Face Recognition


Coding = build a template from an image

- A chain of forward operation in the Network
- Coding is done on CPU or GPU. On a PC, a Smartphone or in the Cloud



Matching = compare 2 templates

- A simple Euclidian distance is used
- Very fast. Millions per second per core






Face Recognition Performances

- Median value on 40 different internal databases.
- Absolute performance varies with images quality

performance evolution





Vendor Benchmark

Confidential / Restricted / Public
Presentation or part title

	Ranked 1st - Minutiae Interoperability Exchange (MINEX III, Ongoing MINEX)
	Ranked 1st - NIST Proprietary Fingerprint Template Test Phase II (PFT II)
	Ranked 1st - NIST Evaluation of Latent Fingerprint Technologies: Extended Feature Sets (ELFT-EFS Image Only)
	Ranked 2nd - Fingerprint Vendor Technology Evaluation 2012 (FPVTE)
	Ranked 1st - Face Recognition Vendor Tests (FRVT), December 2017
	Ranked 2nd – Face in Video Evaluation (FIVE)
	Ranked 1st – Iris Recognition Vendor Test (IREX IV)
	Ranked 1st - Tattoo Recognition Technology – Challenge (Tatt-C 2015)

04/02/2020




DATA



3



HOW TO ACHIEVE?




What people think AI is about



The reality



DATA





Data and label for this session



Face images coming from the Web :

<http://www.cbsr.ia.ac.cn/english/CASIA-WebFace-Database.html>

♦ Dong Yi, Zhen Lei, Shengcai Liao and Stan Z. Li, "Learning Face Representation from Scratch". arXiv preprint arXiv:1411.7923. 2014.



Pre-processed :

centred on the face, 48*48 pixels, grey levels.



Labels :

gender automatically deduced from first name



4 databases :

training : 1000, 10k and 100k

testing : 10k



Goal :

finding Gender from Image





Learning FRAMEWORKs


4



WHY NOW?

Three major recent evolutions


Large Data Set



Computation Power

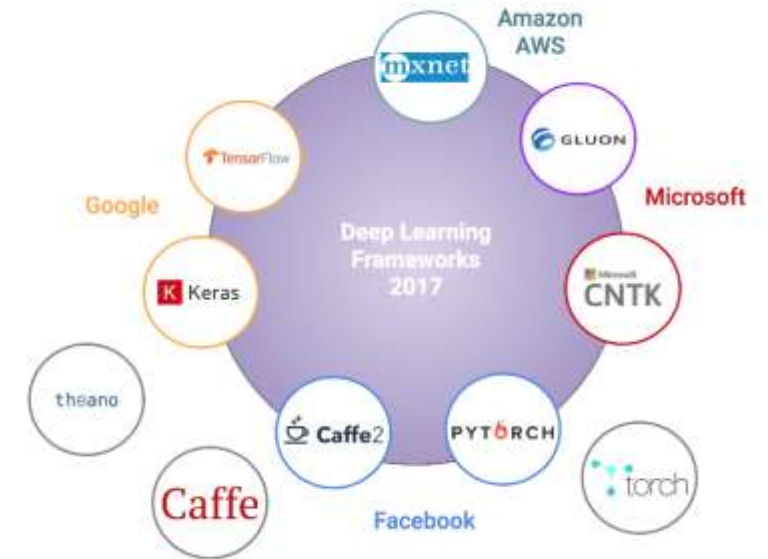


47X Higher Throughput than CPU Server on Deep Learning Inference



Workload: ResNet-50 | CPU: 1X Xeon E5-2650v4 @ 2.5GHz | GPU: add 1X NVIDIA® Tesla® P100 or V100

Deep Learning Frameworks available in OpenSource





Deep Learning Frameworks



Caffe2



theano



TensorFlow



Caffe




Deep learning software

cuDNN is a highly optimized GPU library for NVIDIA units that allows deep learning networks to be trained more quickly

It can dramatically accelerate the performance of a deep network and is often called by the other packages above.

TensorFlow CNN benchmarks

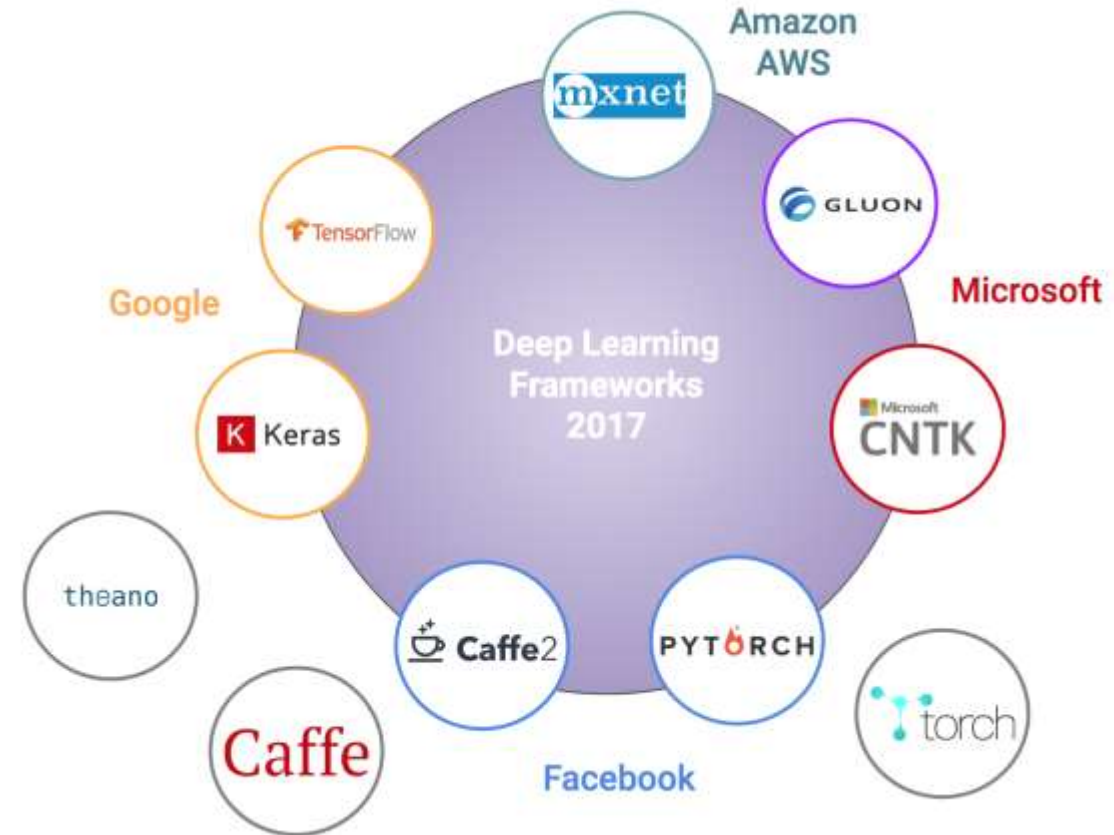




Deep Learning High Level Frameworks

Keras is a Python library that runs on top of either Theano or TensorFlow that allows one to quickly define a network architecture in terms of layers and also includes functionality for image and text preprocessing

Gluon is a framework built on top of MxNet or CNTK. It allows dynamic computation graph.





Neural NETWORK




5




Neural Network: A Neuron

A neuron is a computational unit in the neural network that exchanges messages with each other.




$$a_i \leftarrow g(in_i) = g(\sum_j W_{j,i} a_j)$$






Neural Networks






A deep neural network

Simple Neural Network



● Input Layer

Deep Learning Neural Network





● Hidden Layer

● Output Layer




TRAINING PROCESS



Deep Learning Neural Network



Pixels



Input Layer

Hidden Layer

Output Layer

error

+0.89

-0.02

-0.35

-0.26

Labels

1.0 car

0.0 truck

0.0 airplane

0.0 ship



TensorFlow Playground

<http://playground.tensorflow.org/>

Tinker With a **Neural Network** Right Here in Your Browser.
Don't Worry, You Can't Break It. We Promise.



DATA

Which dataset do you want to use?



Ratio of training to test data: 50%

Noise: 0

Batch size: 10

INPUT

Which properties do you want to feed in?

X_1

X_2

X_1^2

X_2^2

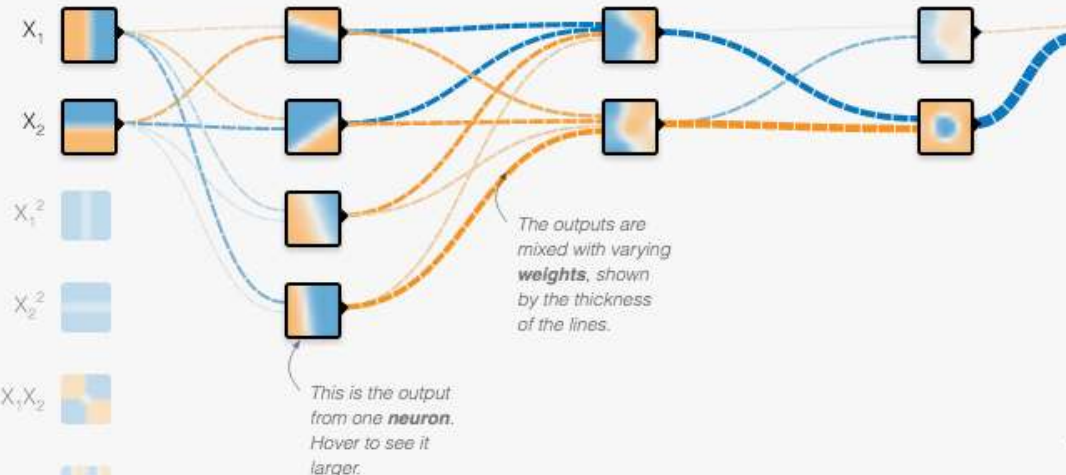
$X_1 X_2$

3 HIDDEN LAYERS

4 neurons


2 neurons

2 neurons



OUTPUT

Test loss 0.000
Training loss 0.000





Tensor Flow



5





Big picture

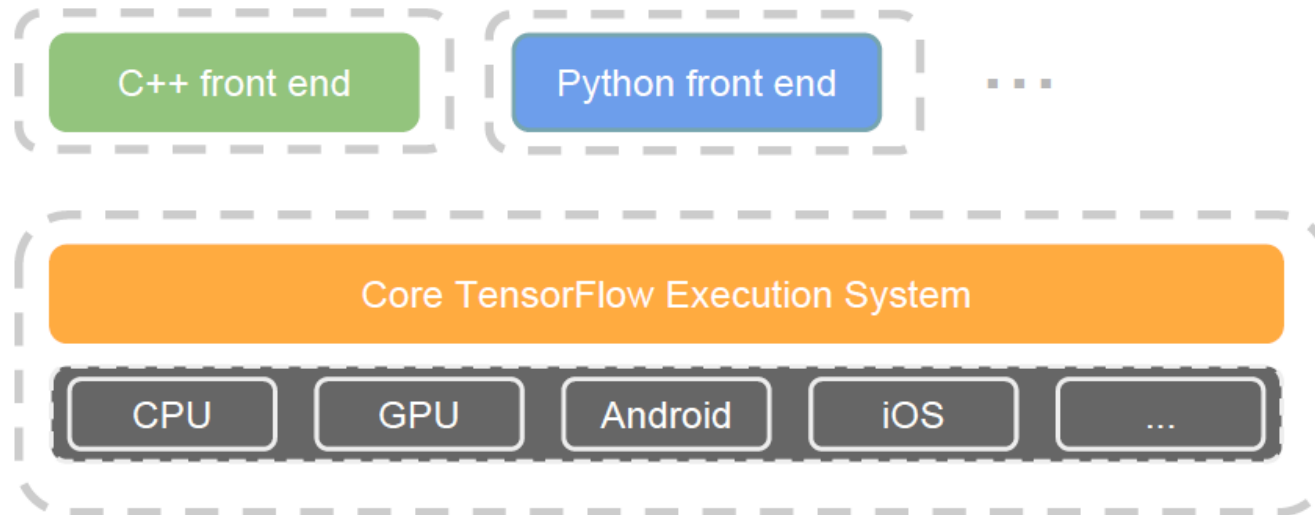
TensorFlow is a deep learning framework released by Google in November 2015

TF is open source

Software architecture

Core in C++

Front ends in Python and C++



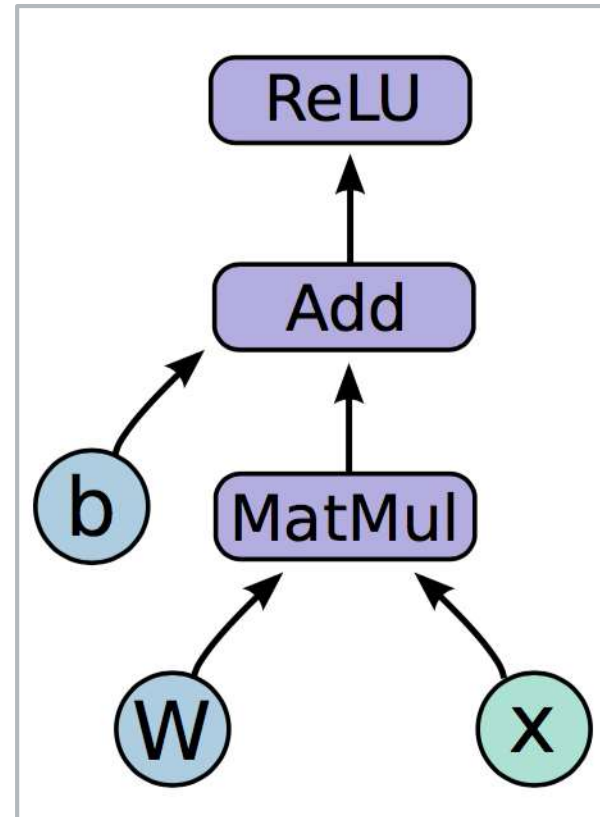


computational graph

Graph nodes are operations which have any number of inputs and outputs

Graph edges are tensors (n-dimensional arrays) which flow between nodes

$$h_i = \text{ReLU}(Wx + b)$$





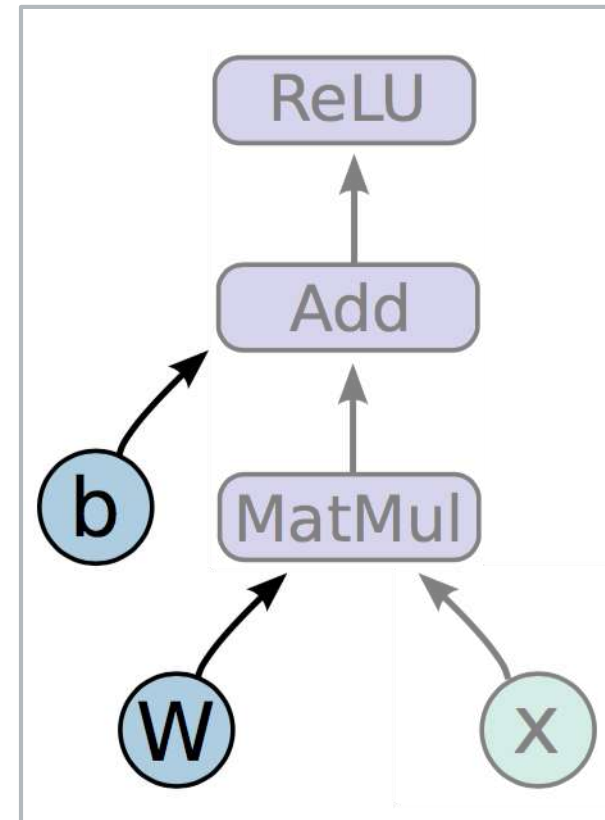
computational graph

Graph nodes are operations which have any number of inputs and outputs

Graph edges are tensors (n-dimensional arrays) which flow between nodes

$$h_i = \text{ReLU}(Wx + b)$$

Variables are nodes which output their current value.
(Parameters to be learned)





Computational graph

Graph nodes are operations which have any number of inputs and outputs

Graph edges are tensors (n-dimensional arrays) which flow between nodes


$$h_i = \text{ReLU}(Wx + b)$$

Mathematical operations:

MatMul: Multiply two matrix values.

Add elementwise (with broadcasting).

ReLU: Activate with elementwise rectified linear function.





A large variety of operators is available

Elementwise math: cos, exp, log, erf, modulo, abs, sqrt, pow, ceil, floor ...

Linear algebra: matrix multiplication, inverse, Cholesky, determinant...

Logical operators: and, or, xor

Comparison operators: >, <, =...

Reduction operators: min, max, argmin, argmax, mean, sum, product...

Slicing/joining: operators: split, concatenate, tile, pad...

Image operators: read/write jpg/png, crop, resize, flip, flop, brightness/contrast/hue adjustment...

Neural network: convolution, pooling, ReLU, sigmoid, dropout, batch normalization, local normalization



Python front end

Everything is writing python code, no prototype files or anything

1. Build python functions describing a graph

Graph contains parameter specifications

Graph contains model architecture

2. Define optimization process

Define values you want to optimize

Attach an optimizer

3. Fetch and feed data

...

You write your optimization loop in python

You can execute whatever python code want to in it!



Monitoring

TensorFlow allows you to monitor pretty much everything in real-time

Values,

Histograms,

Proportion of non-zero values,

Images.

You just have to

Start a web server on the computer where TensorFlow is running

Access to a web page from your own computer (with Chrome)

TensorBoard




Gender Classification with TensorFlow and Python

- **nn1**
the simplest script
- **nn2**
logs in tensorboard
- **nn3**
Compute accuracy on a test database
- **nn4**
use convolutional network
- **nn5**
use dropout
- **From Session2 to Session3**
your turn => learn a CNN and fool it !



nn1 : A simple perceptron





nn1 : simplest.py

```
import tensorflow as tf
import numpy as np

# nombre d images
nbdata = 1000
trainDataFile = '../DataBases/data_1k.bin'
LabelFile = '../DataBases/gender_1k.bin'

# taille des images 48*48 pixels en niveau de gris
dim = 2304
f = open(trainDataFile, 'rb')
data = np.empty([nbdata, dim], dtype=np.float32)
for i in range(nbdata):
    data[i, :] = np.fromfile(f, dtype=np.uint8,
count=dim).astype(np.float32)
f.close()

f = open(LabelFile, 'rb')
label = np.empty([nbdata, 2], dtype=np.float32)
for i in range(nbdata):
    label[i, :] = np.fromfile(f, dtype=np.float32, count=2)
f.close()

class fc_layer(tf.Module):
    def __init__(self, input_dim, output_dim):
        w_init = tf.random.truncated_normal([input_dim, output_dim],
stddev=0.1)
        self.w = tf.Variable(w_init)
        print('w', self.w.get_shape())
        b_init = tf.constant(0.0, shape=[output_dim])
        self.b = tf.Variable(b_init)
        print('b', self.b.get_shape())

    def __call__(self, x):
        return tf.matmul(x, self.w) + self.b
```

```
class SimpleNet(tf.Module):
    def __init__(self, input_dim):
        self.fc1 = fc_layer(input_dim, 50)
        self.fc2 = fc_layer(50, 2)

    def __call__(self, x):
        x = self.fc1(x)
        x = tf.nn.sigmoid(x)
        x = self.fc2(x)
        return x

def train_one_step(model, optimizer, image, label):
    with tf.GradientTape() as tape:
        y = model(image)
        loss = tf.reduce_sum(tf.square(y - label))
        grads = tape.gradient(loss, model.trainable_variables)
        optimizer.apply_gradients(zip(grads, model.trainable_variables))
    return loss

optimizer = tf.optimizers.SGD(1e-5)
simple_model = SimpleNet(dim)

curPos = 0
batchSize = 256

for it in range(5000):
    if curPos + batchSize > nbdata:
        curPos = 0
    loss = train_one_step(simple_model, optimizer,
data[curPos:curPos + batchSize, :],
label[curPos:curPos + batchSize, :])

    curPos += batchSize
    if it % 100 == 0 or it < 10:
        print("it= %6d - loss= %f" % (it, loss.numpy()))
```



nn1 : let's run

```
urld27@Deep6: ~/Partage/Stephane/td/nn1
18:43:07|nn1> python simplest.py
I tensorflow/stream_executor/dso_loader.cc:108] successfully opened CUDA library libcublas.so locally
I tensorflow/stream_executor/dso_loader.cc:108] successfully opened CUDA library libcudnn.so locally
I tensorflow/stream_executor/dso_loader.cc:108] successfully opened CUDA library libcufft.so locally
I tensorflow/stream_executor/dso_loader.cc:108] successfully opened CUDA library libcuda.so.1 locally
I tensorflow/stream_executor/dso_loader.cc:108] successfully opened CUDA library libcurand.so locally
Tensor("truncated_normal:0", shape=(2304, 50), dtype=float32)
<tensorflow.python.ops.variables.Variable object at 0x7f7b7ce0f310>
Tensor("truncated_normal_1:0", shape=(50, 2), dtype=float32)
<tensorflow.python.ops.variables.Variable object at 0x7f7b586d38d0>
I tensorflow/core/common_runtime/gpu/gpu_init.cc:102] Found device 0 with properties:
name: GeForce GTX 1080
major: 6 minor: 1 memoryClockRate (GHz) 1.7335
pciBusID 0000:02:00:0
Total memory: 7.92GiB
Free memory: 7.81GiB
I tensorflow/core/common_runtime/gpu/gpu_init.cc:126] DMA: 0
I tensorflow/core/common_runtime/gpu/gpu_init.cc:136] 0: Y
I tensorflow/core/common_runtime/gpu/gpu_device.cc:838] Creating TensorFlow device (/gpu:0) -> (device: 0, name: GeForce GTX 1080, pci bus id: 0000:02:00:0)
it= 0 - loss= 181.141205
it= 1000 - loss= 35.203453
it= 2000 - loss= 39.134914
it= 3000 - loss= 53.792587
it= 4000 - loss= 23.412951
it= 5000 - loss= 30.071821
it= 6000 - loss= 47.140503
it= 7000 - loss= 20.690556
it= 8000 - loss= 26.863918
it= 9000 - loss= 45.297955
it= 10000 - loss= 21.248180
it= 11000 - loss= 26.216721
it= 12000 - loss= 42.632809
it= 13000 - loss= 18.814625
it= 14000 - loss= 23.940264
it= 15000 - loss= 42.456089
```



nn1 : simplest.py

What are the shapes of these tensors ?

```
class fc_layer(tf.Module):  
    def __init__(self, input_dim, output_dim):  
        w_init = tf.random.truncated_normal([input_dim, output_dim], stddev=0.1)  
        self.w = tf.Variable(w_init)  
        print('w', self.w.get_shape())  
        b_init = tf.constant(0.0, shape=[output_dim])  
        self.b = tf.Variable(b_init)  
        print('b', self.b.get_shape())  
  
    def __call__(self, x):  
        return tf.matmul(x, self.w) + self.b
```

How many parameters to learn ?



nn1 : simplest.py

```
nn2 ConvNeuralNet > __call__()
```

```
Run: simplest x
```

```
TensorFlow version: 2.1.0
2020-02-02 19:50:34.939551: W tensorflow/stream_executor
2020-02-02 19:50:34.940421: E tensorflow/stream_executor
2020-02-02 19:50:34.966690: I tensorflow/stream_executor
2020-02-02 19:50:34.968776: I tensorflow/stream_executor
2020-02-02 19:50:34.969818: I tensorflow/core/platform
w      (2304, 50)
b      (50,)
w      (50, 2)
b      (2,)
2020-02-02 19:50:35.010644: W tensorflow/python/util/util.cc:319] Sets are not currently considered
them.
it=    0 - loss= 357.267761
it=    1 - loss= 112.971466
it=    2 - loss= 98.500824
it=    3 - loss= 303.063049
it=    4 - loss= 91.577881
it=    5 - loss= 87.436913
it=    6 - loss= 289.197998
it=    7 - loss= 83.190666
it=    8 - loss= 81.066132
it=    9 - loss= 282.216309
it=   100 - loss= 76.240219
it=   200 - loss= 70.685059
it=   300 - loss= 214.504547
```

4: Run 6: TODO Terminal Python Console

What are the shapes of these tensors ?

$2304 * 50$

50

$50 * 2$

2

How many parameters to learn ?

115352



nn2 : simplest_v2.py

Save and Load the network variables

```
ckpt = tf.train.Checkpoint(step, optimizer, net)
ckpt.save('filename')
ckpt.restore('filename-step')
```

Following the learning process : Summary

```
t = tf.summary.create_file_writer('logs')
```

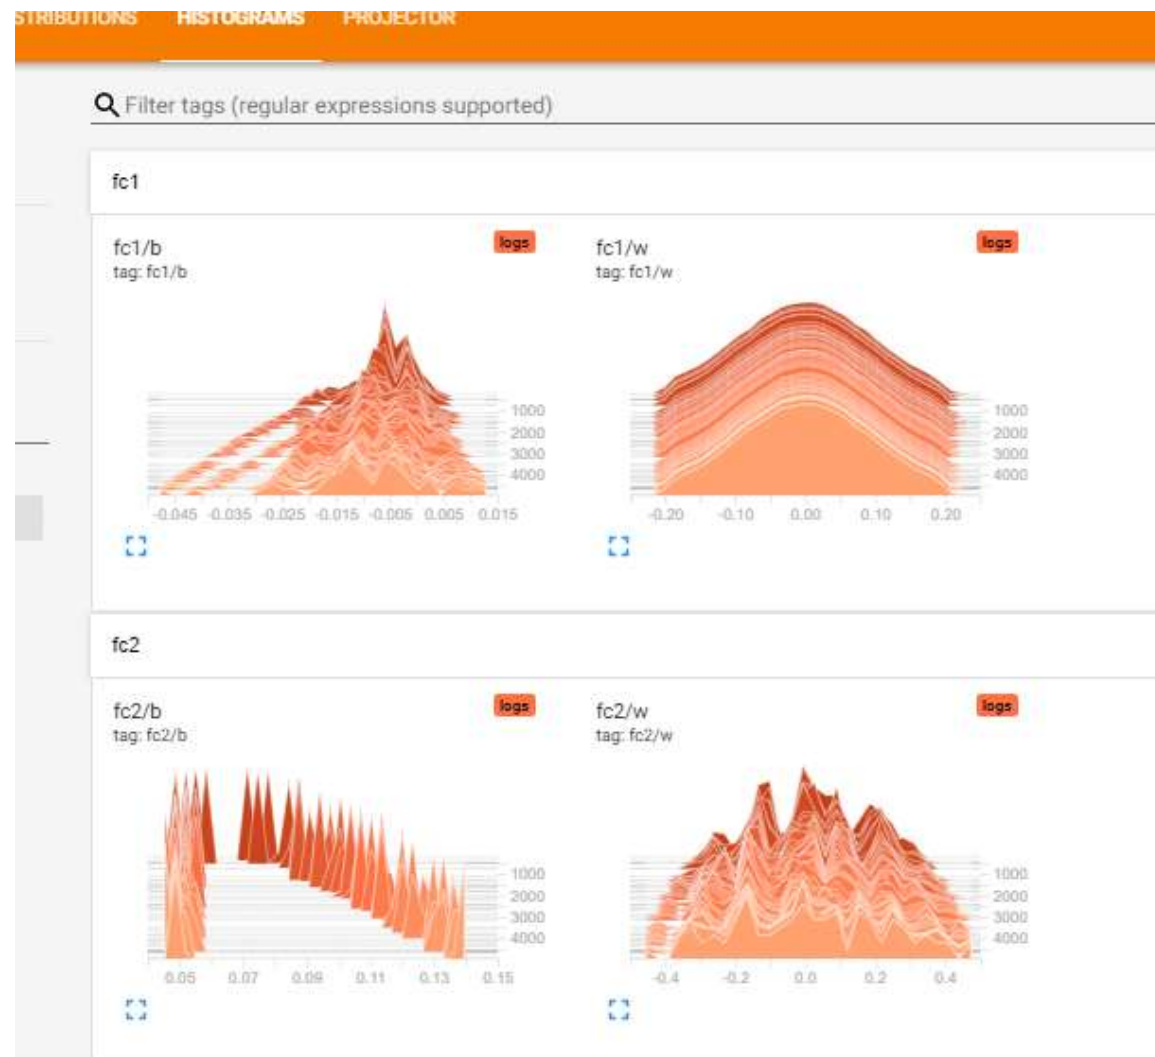
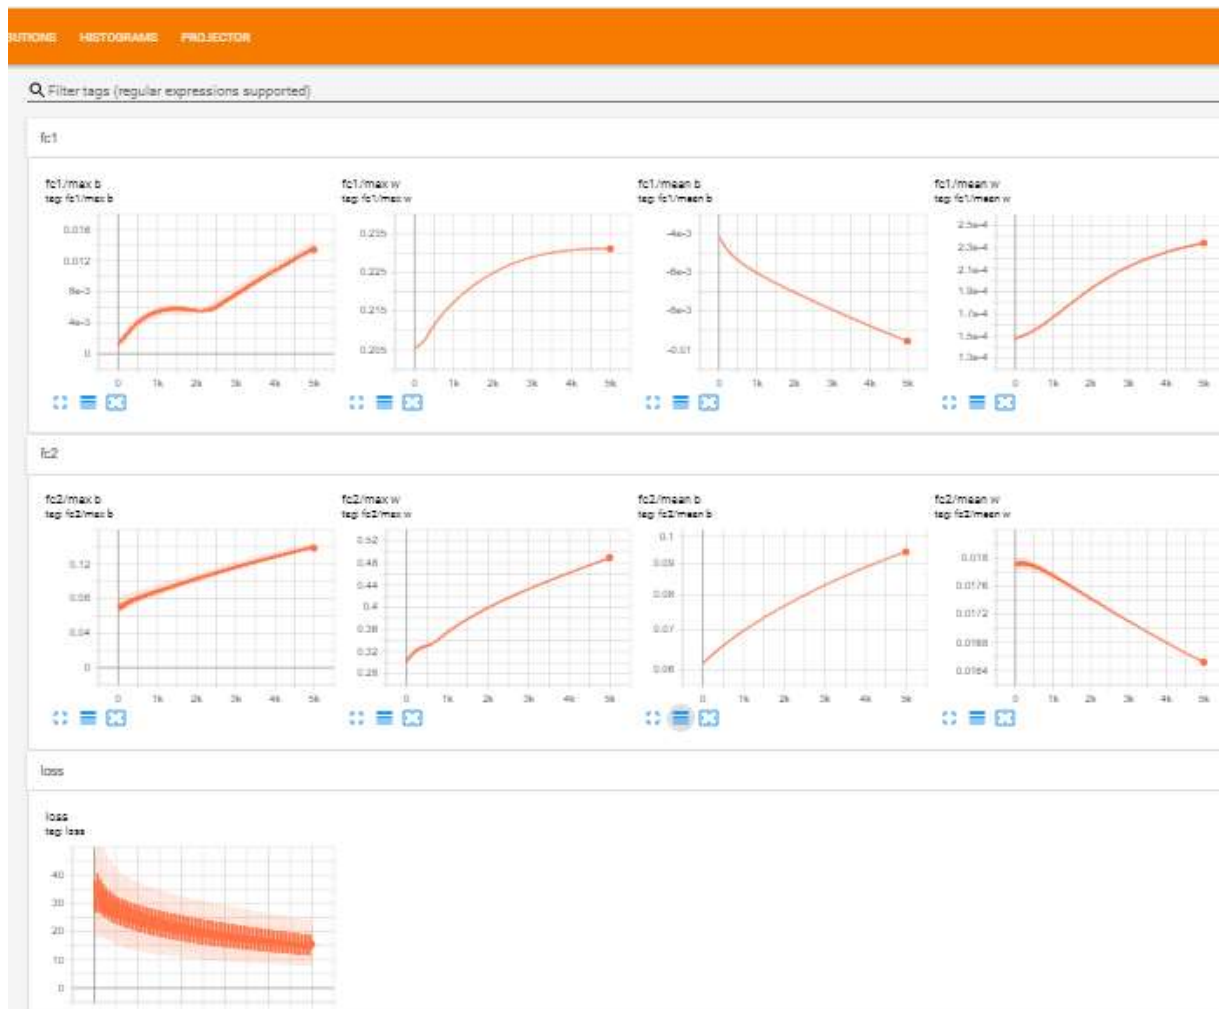
```
with t.as_default():
```

```
    tf.summary.scalar(name, tensor)
    tf.summary.histogram(name, tensor)
```



nn2 : simplest_v2.py

tensorboard --logdir nn2 --port 8888





nn3 : perceptron for gender classification ... does it work ?

3 datasets for learning and 1 for testing

```
train = ds.DataSet('../DataBases/data_1k.bin', '../DataBases/gender_1k.bin', 1000)
train = ds.DataSet('../DataBases/data_10k.bin', '../DataBases/gender_10k.bin', 10000)
train = ds.DataSet('../DataBases/data_100k.bin', '../DataBases/gender_100k.bin', 100000)
test = ds.DataSet('../DataBases/data_test10k.bin', '../DataBases/gender_test10k.bin', 10000)
```

Compute Accuracy

```
correct_prediction = tf.equal(tf.argmax(y, 1), tf.argmax(labels, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
```

On a complete dataset

```
def mean_accuracy(self, model):
    acc = 0
    for i in range(0, self.npdata, self.batchSize):
        curBatchSize = min(self.batchSize, self.npdata - i)
        images = self.data[i:i+curBatchSize,:]
        labels = self.label[i:i+curBatchSize,:]
        y = model(images, False)
        correct_prediction = tf.equal(tf.argmax(y, 1), tf.argmax(labels, 1))
        accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
        acc += accuracy * curBatchSize
    acc /= self.npdata
    tf.summary.scalar('Accuracy %s'%self.name, acc)
    return acc
```



CNN Architecture

6

- Fully Connected layer
- Activation layer
- Convolutional Layer
- Pooling Layer
- Loss Layer



Fully Connected layer



6.1



Fully Connected Layer



All the neurons of a layer are connected to the next layer.



If X is the input vector of values, Y the output vector and W the parameter matrix.



The computation of Y is a simple scalar product. $Y = W.X$




FC layers (Fully connected) are often named IP layers (Inner Product)





A regular 3-layer Fully Connected Neural Network





Activation Functions

6.2



Sigmoid function

<http://mathworld.wolfram.com/SigmoidFunction.html>

$$f(u) = \frac{1}{1 + e^{-\beta u}} = \frac{1}{1 + e^{-u}}, \text{ for simplicity set } \beta = 1$$


$$\frac{df(u)}{du} = f'(u)$$

Hence

$$f'(u) = \frac{df(u)}{du} = \frac{d\left(\frac{1}{1 + e^{-u}}\right)}{d(1 + e^{-u})} \frac{d(1 + e^{-u})}{du}, \text{ (using chain rule)}$$

$$\begin{aligned} f'(u) &= \frac{-1}{(1 + e^{-u})^2} (-e^{-u}) = \frac{1}{(1 + e^{-u})^2} e^{-u} \\ &= \frac{1}{(1 + e^{-u})} \frac{e^{-u}}{(1 + e^{-u})} = \frac{1}{(1 + e^{-u})} \frac{[(1 + e^{-u}) - (1 + e^{-u})] + e^{-u}}{(1 + e^{-u})} \\ &= \frac{1}{(1 + e^{-u})} \left(1 - \frac{1}{(1 + e^{-u})}\right) = f(u)(1 - f(u)) \end{aligned}$$

$$\text{Thus, } \frac{df(u)}{du} = f'(u) = f(u)(1 - f(u))$$





Rectified Linear Units

$$\text{Relu} (x) = \text{MAX} (0 , x)$$

More efficient gradient propagation, derivative is 0 or constant, just fold into learning rate

More efficient computation: Only comparison, addition and multiplication.

A variation : Leaky ReLU $f(x) = x$ if $x > 0$ else ax where $0 \leq a \leq 1$, so that derivative is not 0 and can do some learning for that case.

Lots of other variations

Sparse activation: For example, in a randomly initialized networks, only about 50% of hidden units are activated (having a non-zero output)



Convolutional Layer

6.3




Background


Convolution is an operation on two functions f and g , which produces a third function that can be interpreted as a modified ("filtered") version of f . In this interpretation we call g the *filter (or Kernel)*

$$f(x) * g(x) = \int_{-\infty}^{\infty} f(\tau) \cdot g(x - \tau) d\tau$$


Convolution of a block with a block



Convolution of a triangle with a block



edge detection





Background

For functions of a discrete variable x , i.e. arrays of numbers, the definition is:

$$f[x] * g[x] = \sum_{k=-\infty}^{\infty} f[k] \cdot g[x - k]$$

Finally, for functions of two variables x and y (for example images), these definitions become:

$$f(x,y) * g(x,y) = \int_{\tau_1=-\infty}^{\infty} \int_{\tau_2=-\infty}^{\infty} f(\tau_1, \tau_2) \cdot g(x - \tau_1, y - \tau_2) d\tau_1 d\tau_2$$

And :

$$f[x,y] * g[x,y] = \sum_{n_1=-\infty}^{\infty} \sum_{n_2=-\infty}^{\infty} f[n_1, n_2] \cdot g[x - n_1, y - n_2]$$

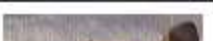

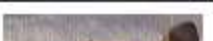

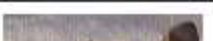

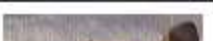



Background

Convolutional Neural Networks

Variation of multi-layer neural networks

Kernel (Convolution Matrix)

Identity	Edge detection	[1 0 -1]		
Sharpen	Edge detection	[1 0 -1]		
Box blur (normalized)	Edge detection	[1 0 -1]		
Gaussian blur (approximation)	Edge detection	[1 0 -1]		

reference : [http://en.wikipedia.org/wiki/Kernel_\(image_processing\)](http://en.wikipedia.org/wiki/Kernel_(image_processing))



Background

Convolutional Filter



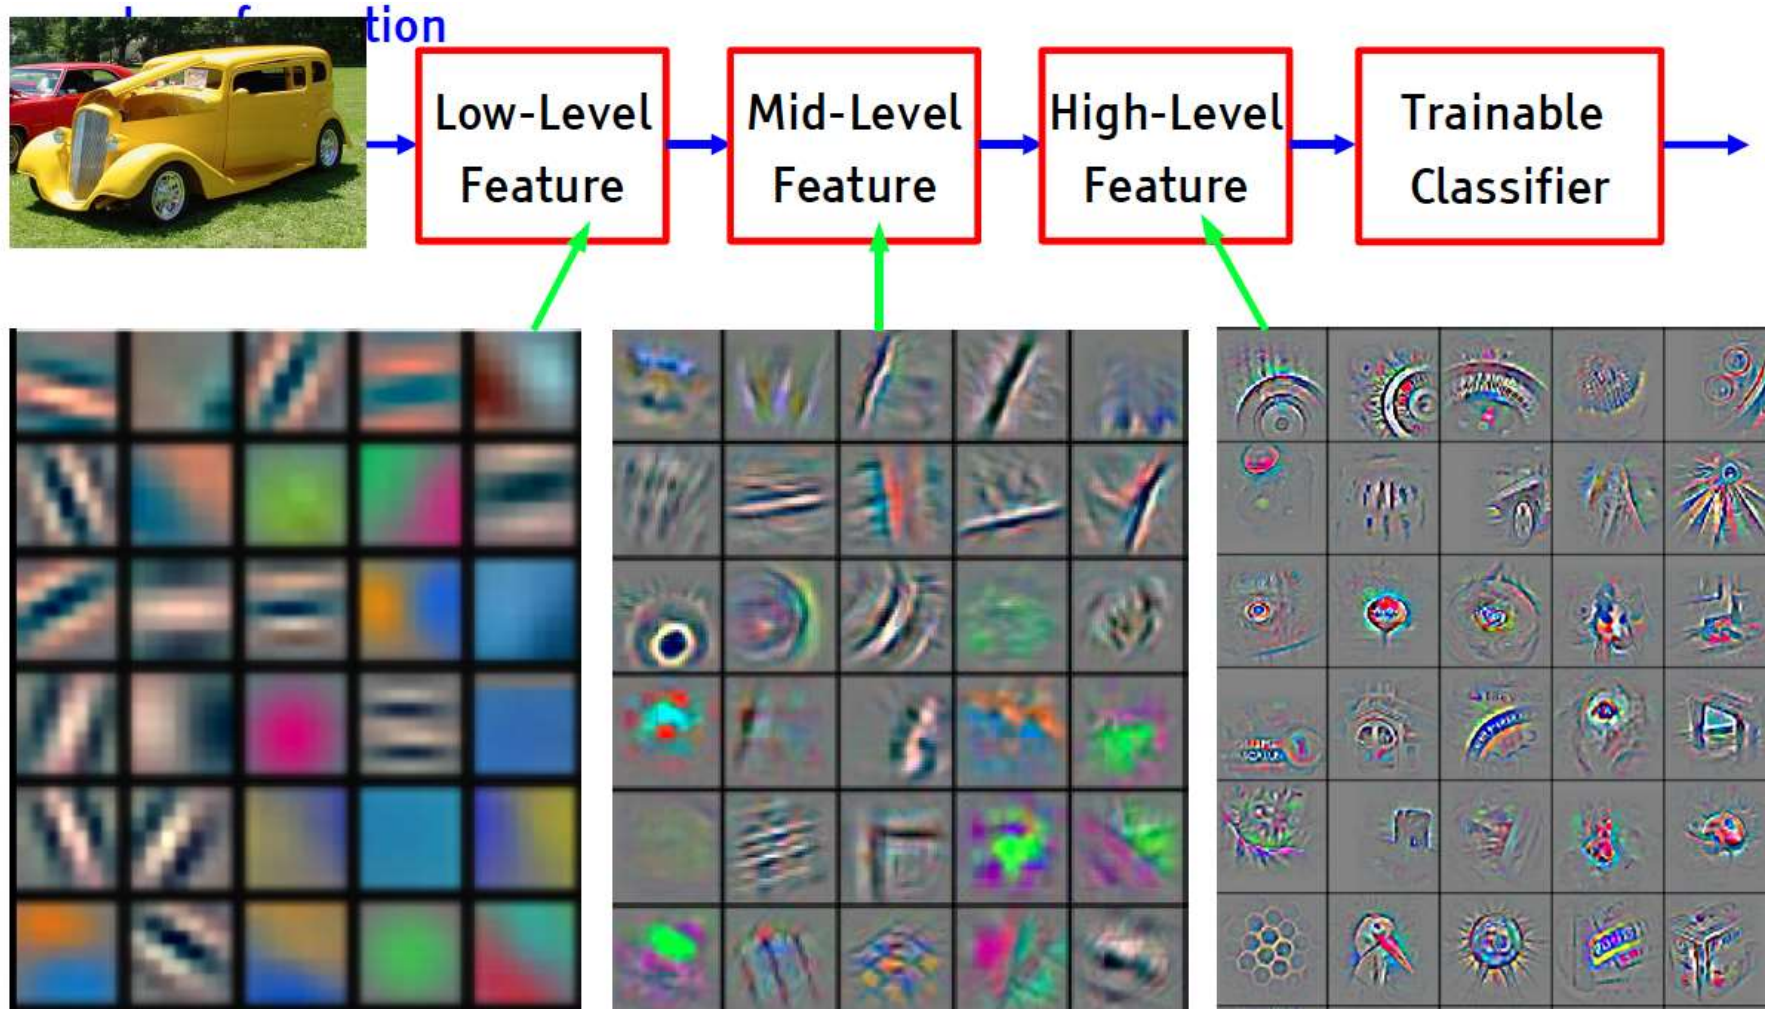
Input



Feature Map




Deep Learning = Learning Hierarchical Representations



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]




A ConvNet looks at data in three dimensions (width, height, depth)






The activations of an example ConvNet architecture.





ConvNets



Input : “InDim” images with size H,W

A Convolution Layer : “outDim” filters with kernel size InDim x KsizeH x KsizeW.

Output : “outDim” images with size H,W (if stride == 1)



pooling Layer




6.4





Pooling






Pooling


Spatial Pooling

Non-overlapping / overlapping regions


Average or Max



Max




Avg





max pooling





Loss Layer



6.5





Minimum Mean Squared Error

Works generally well for learning tasks

$$\mathcal{L} = \sum_n (y^{(n)} - t^{(n)})^2$$

$y^{(n)}$ - output of the classifier

$t^{(n)}$ - labels

$$y^{(n)} = F(\mathbf{x}^{(n)})$$
$$F \in \mathcal{F}$$

\mathcal{F} - the set of all functions representable by the neural network architecture.

$$\min_{F \in \mathcal{F}} \sum_n (F(\mathbf{x}^{(n)}) - t^{(n)})^2$$



Softmax function

Definition

$$\mathbf{y}^{(n)}(i) = \frac{\exp\left(-\mathbf{a}^{(n)}(i)\right)}{\sum_j \exp\left(-\mathbf{a}^{(n)}(j)\right)}$$

It can be easily verify that

$$\sum_i \mathbf{y}^{(n)}(i) = 1$$



Cross Entropy

Definition: $p_X(i), q_X(i)$

$$H(p_X|q_X) = - \sum_i p_X(i) \log q_X(i)$$

The cross entropy loss is closely related to the Kullback-Leibler divergence between the empirical distribution and the predicted distribution.

Cross Entropy as a Loss Function:

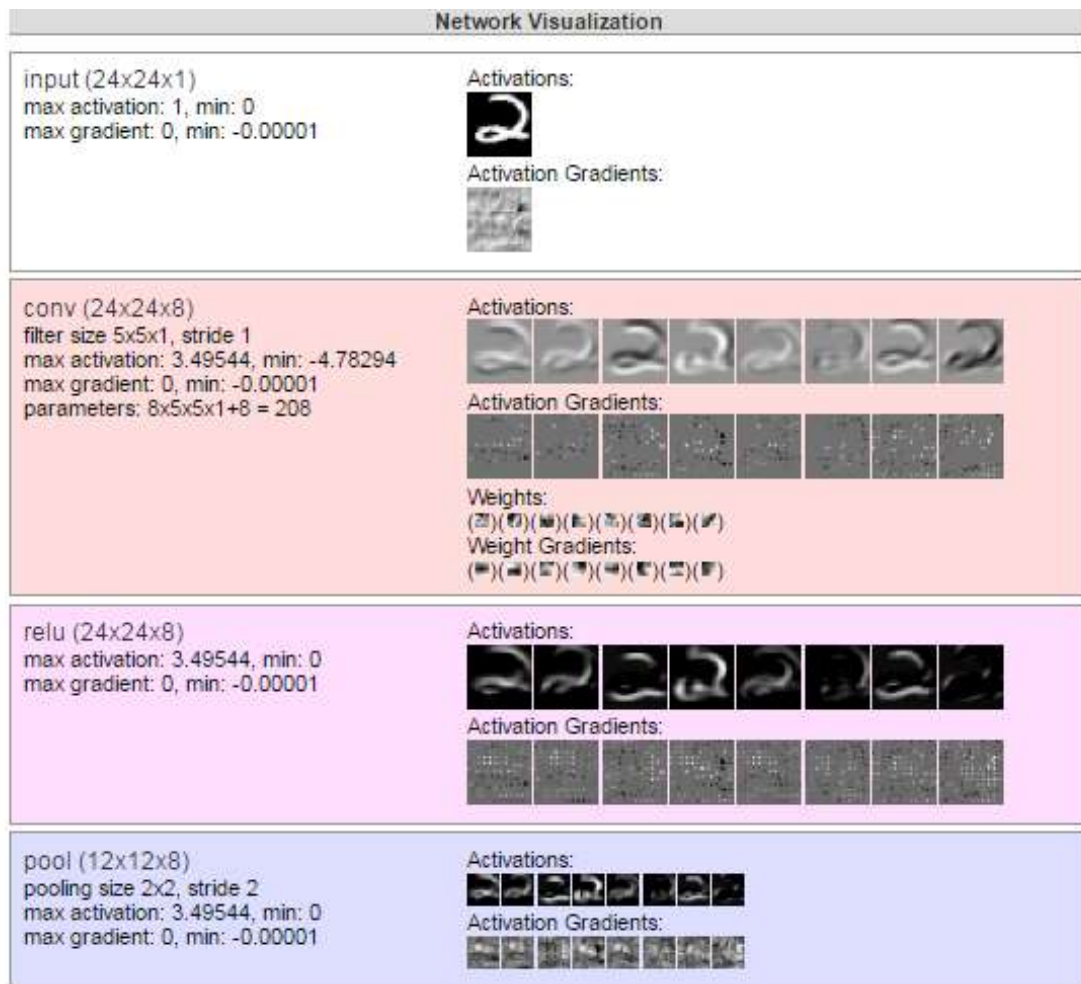
$$\ell^{(n)} = H(t^{(n)}|y^{(n)}) = - \sum_i t^{(n)}(i) \log y^{(n)}(i)$$

$$\mathcal{L} = \sum_n \ell^{(n)} = - \sum_n \sum_i t^{(n)}(i) \log y^{(n)}(i)$$



Another CNN visualization

<https://cs.stanford.edu/people/karpathy/convnetjs/demo/mnist.html>



conv (12x12x16)
filter size 5x5x8, stride 1
max activation: 8.19594, min: -14.80998
max gradient: 0, min: -0.00001
parameters: $16 \times 5 \times 5 \times 8 + 16 = 3216$

Activations:



Activation Gradients:



Weights:



Weight Gradients:



relu (12x12x16)
max activation: 8.19594, min: 0
max gradient: 0, min: -0.00001

Activations:



Activation Gradients:



pool (4x4x16)
pooling size 3x3, stride 3
max activation: 8.19594, min: 0
max gradient: 0, min: -0.00001

Activations:



Activation Gradients:



fc (1x1x10)
max activation: 7.62757, min: -19.6438
max gradient: 0, min: -0.00001
parameters: $10 \times 256 + 10 = 2570$

Activations:




Activation Gradients:



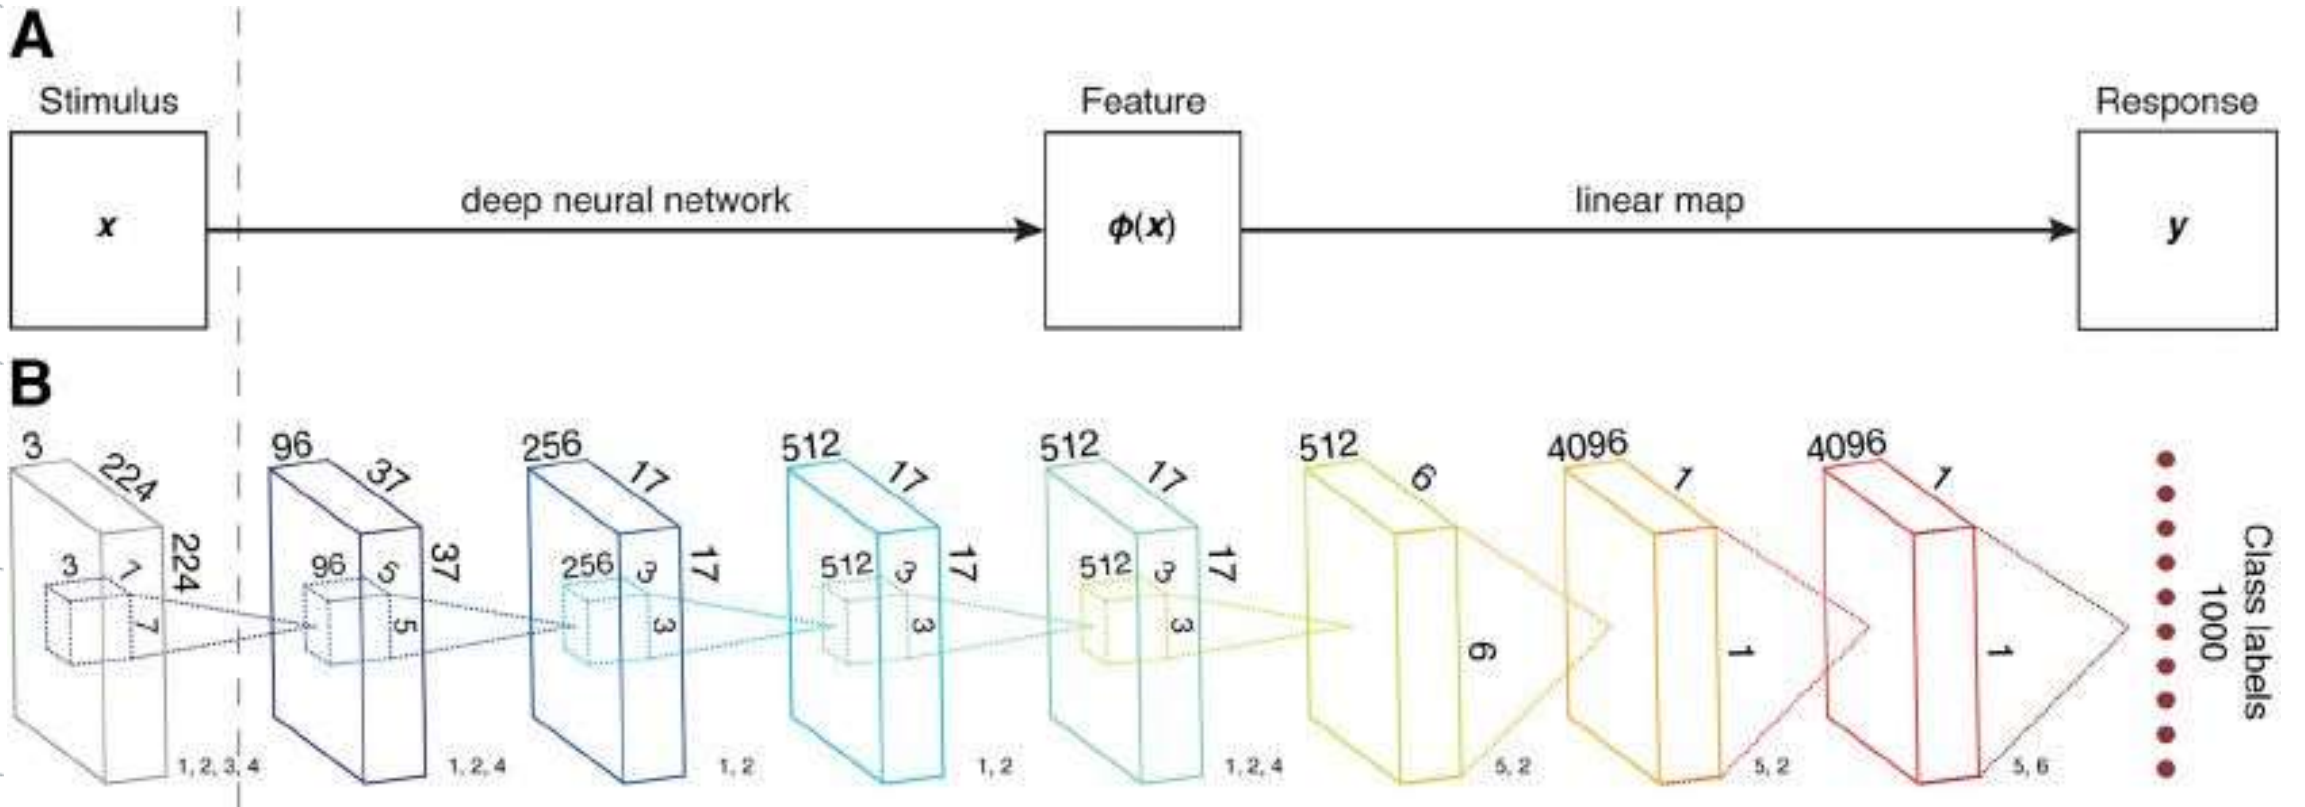
softmax (1x1x10)
max activation: 0.99999, min: 0
max gradient: 0, min: 0

Activations:





Example of deep convolutional neural network for image classification



ImageNet : 12 M images

Networks : up to 60 M parameters



CNN Building Blocks

- **Fully Connected layer**

`tensor = tf.matmul(tensor, W) + B`

- **Activation layer**

`tensor = tf.nn.relu(tensor)`

- **Convolutional Layers**

`tensor = tf.nn.conv2d(tensor, W, strides, padding) + B`

- **Pooling Layers**

`tensor = tf.nn.max_pool(tensor, ksize, strides, padding)`

- **Loss Layer**

`tensor = tf.nn.log_softmax(tensor)`



nn4 : a CNN

Definition of the network :

```
class ConvNeuralNet(tf.Module):
    def __init__(self):
        self.unflat = Layers.unflat('unflat', 48, 48, 1)
        self.cv1 = Layers.conv('conv_1', output_dim=3, filterSize=3, stride=1)
        self.mp = Layers.maxpool('pool', 2)
        self.cv2 = Layers.conv('conv_2', output_dim=6, filterSize=3, stride=1)
        self.cv3 = Layers.conv('conv_3', output_dim=12, filterSize=3, stride=1)
        self.flat = Layers.flat()
        self.fc = Layers.fc('fc', 2)

    def __call__(self, x, log_summary):
        x = self.unflat(x, log_summary)
        x = self.cv1(x, log_summary)
        x = self.mp(x)
        x = self.cv2(x, log_summary)
        x = self.mp(x)
        x = self.cv3(x, log_summary)
        x = self.mp(x)
        x = self.flat(x)
        x = self.fc(x, log_summary)
        return x
```



nn4 : number of parameters and shapes ...

Total number of parameters learned in nn4 ?

$$\text{Conv 1} \Rightarrow (3*3*1+1)*3 = 30$$

$$\text{Conv 3} \Rightarrow (3*3*3+1)*6 = 168$$

$$\text{Conv 5} \Rightarrow (3*3*6+1)*12 = 660$$

$$\text{FC 7} \Rightarrow (6*6*12+1)*2 = 866$$

$$\text{Total} = 1724$$

Tensor shape as input for Conv_3 ?

In Dataset.py, default batchsize = 128

Shape = (128,24,24,3)