

Reinforcement Learning: Methods and Algorithms


Thomas Bonald

2019 – 2020




Reinforcement Learning


- ▶ Learning by **trial and error**
- ▶ Inspired by the behavior of animals (including humans!)
- ▶ The **exploration-exploitation** trade-off
- ▶ Many applications: robotics, games, advertising, content recommendation, medicine, etc.



Toy example: Maze



Toy example: Pac-Man



Outline

1. **Markov decision process**
2. Dynamic programming
3. Online estimation
4. Online control

Markov decision process

At time $t = 0, 1, 2, \dots$, the agent in **state** s_t takes **action** a_t and:

- ▶ receives **reward** r_t
- ▶ moves to **state** s_{t+1}

The reward and new state are **stochastic** in general.

Some states may be **terminal**.

Definition

A **Markov decision process** (MDP) is defined by:

- ▶ some initial state s_0
- ▶ the reward distribution, $r_t \sim p(r|s_t, a_t)$
- ▶ the transition probabilities, $s_{t+1} \sim p(s|s_t, a_t)$

Policy

Definition

Given a Markov decision process, a **policy** defines the action taken in each state:

$$\pi(a|s) = P(a_t = a \mid s_t = s)$$

- ▶ A policy is **stochastic** in general.
- ▶ When **deterministic**, we use the simple notation $\pi(s)$ for the action taken in state s .

Remark

Given some policy π , the sequence of states s_0, s_1, s_2, \dots defines a **Markov process**.

Objective function

Definition

Given the rewards r_0, r_1, r_2, \dots , we refer to the **gain** as:

$$\begin{aligned} G &= r_0 + \gamma r_1 + \gamma^2 r_2 + \dots \\ &= \sum_{t=0}^{+\infty} \gamma^t r_t \end{aligned}$$

The parameter $\gamma \in [0, 1]$ is the **discount factor**:

- ▶ $\gamma = 0 \longrightarrow$ immediate reward
- ▶ $\gamma = 1 \longrightarrow$ cumulative reward


Value function

Definition


The **value** function is the expected gain from each state:

$$\forall s, \quad V_{\pi}(s) = \mathbb{E}_{\pi}(G | s_0 = s)$$

Toy example: Maze (random policy)



Toy example: Pac-Man (random policy)



Bellman's equation

Recall the definition:

$$\forall s, \quad V_{\pi}(s) = \mathbb{E}_{\pi}(G \mid s_0 = s)$$

Proposition

The value function V_{π} of any policy π is the unique solution to the **fixed-point equation**:

$$\forall s, \quad V(s) = \mathbb{E}_{\pi}(r_0 + \gamma V(s_1) \mid s_0 = s)$$

Solution to Bellman's equation

Let T_π the **operator** associated with Bellman's equation:

$$\forall s, \quad T_\pi(V)(s) = E_\pi(r_0 + \gamma V(s_1) | s_0 = s)$$

This operator is **contracting**:

$$\forall U, V, \quad \|T_\pi(V) - T_\pi(U)\|_\infty \leq \gamma \|V - U\|_\infty$$

Proposition

If $\gamma < 1$, then:

$$\lim_{n \rightarrow +\infty} T_\pi^n(V) = V_\pi$$

Optimal policy

Definition

A policy π^* is **optimal** if and only if

$$\forall s, \quad V_{\pi^*}(s) \geq V_{\pi}(s)$$

Bellman's optimality equation

Proposition

There is a unique solution V_* to the **fixed-point equation**:

$$\forall s, \quad V(s) = \max_a \mathbb{E}(r_0 + \gamma V(s_1) | s_0 = s, a_0 = a)$$

Optimal value evaluation

Let T_\star the operator associated with Bellman's optimality equation:

$$\forall s, \quad T_\star(V)(s) = \max_a E(r_0 + \gamma V(s_1) | s_0 = s, a_0 = a)$$

Proposition

If $\gamma < 1$, then:


$$\lim_{n \rightarrow +\infty} T_\star^n(V) = V_\star \geq \max_{\pi} V_\pi$$

The proof relies on the fact that:


- ▶ T_\star is **contracting**
- ▶ $T_\star \geq T_\pi$ for any policy π , so that:

$$V_\star = \lim_{n \rightarrow +\infty} T_\star^n(V) \geq \lim_{n \rightarrow +\infty} T_\pi^n(V) = V_\pi$$

Toy example: Maze



Toy example: Pac-Man



Optimal policy

Define:

$$\forall s, \quad \pi^*(s) = a^* \in \arg \max_a \mathbb{E}(r_0 + \gamma V_*(s_1) | s_0 = s, a_0 = a)$$

Bellman's optimality theorem


The policy π^* is optimal:

$$\forall s, \quad V_{\pi^*}(s) = \max_{\pi} V_{\pi}(s)$$


Note that:

- ▶ The policy π^* is **deterministic**
- ▶ The optimal policy is **not unique**.

Toy example: Maze (optimal policy)



Toy example: Pac-Man (optimal policy)



Outline

1. Markov decision process
2. **Dynamic programming** → Policy iteration, Value iteration
3. Online estimation
4. Online control

Policy improvement

Given some policy π , let π' the policy defined by:

$$\pi'(s) = a^* \in \arg \max_a \mathbb{E}(r_0 + \gamma V_\pi(s_1) | s_0 = s, a_0 = a)$$

Proposition

The policy π' is better than π :

$$\forall s, \quad V_{\pi'}(s) \geq V_\pi(s)$$

Policy iteration

Algorithm

Starting from some arbitrary **policy** $\pi = \pi_0$, iterate until convergence:


1. **Evaluate** the policy (by solving Bellman's equation)
2. **Improve** the policy:

$$\forall s, \quad \pi(s) \leftarrow \arg \max_a E(r_0 + \gamma V_\pi(s_1) | s, a)$$


- ▶ The sequence $\pi_0, \pi_1, \pi_2, \dots$ is **monotonic** and converges in **finite time** (for finite number of states and actions).
- ▶ The limit is an **optimal policy**.

Toy example: Maze (policy iteration)


$n = 1$




$n = 2$




$n = 5$



$n = 6$



Toy example: Pac-Man (policy iteration)



Practical considerations

- ▶ The step of **policy evaluation** is time-consuming (solution of Bellman's equation)
- ▶ Do we need the **exact** solution?
No, since it is used only to **improve** the policy!
- ▶ The number of iterations can be limited to some value k
How to set k ? Why not... $k = 1$?

Value Iteration

Algorithm

Starting from some arbitrary **value** function $V = V_0$, iterate until convergence:

$$\forall s, \quad V(s) \leftarrow \max_a E(r_0 + \gamma V(s_1) | s, a)$$

- ▶ The sequence V_0, V_1, V_2, \dots converges in **finite time** (for finite number of states and actions).
- ▶ The **limit** solves Bellman's optimality equation.
- ▶ The corresponding policy is **optimal**.

Outline

1. Markov decision process
2. Dynamic programming
3. **Online estimation** → Monte-Carlo, TD learning
4. Online control

Temporal difference

Assume you want to estimate the **empirical mean** of some data stream x_1, x_2, \dots

Two strategies:

- Store the sum:

For each new sample x_t ,

$$S \leftarrow S + x_t \quad X \leftarrow \frac{S}{t}$$

- Update with the **temporal difference**:

For each new sample x_t ,

$$X \leftarrow X + \alpha(x_t - X) \quad \alpha = \frac{1}{t}$$

Note: Most often, the learning rate α is fixed!

MC learning

Idea: Evaluate the **value** function of some policy π using **episodes** s_0, s_1, \dots, s_T (assuming the presence **terminal** states)

Let:

$$G_0 = r_0 + \gamma r_1 + \dots + \gamma^{T-1} r_{T-1}$$

$$G_1 = r_1 + \gamma r_2 + \dots + \gamma^{T-2} r_{T-1}$$

...

$$G_{T-1} = r_{T-1}$$

MC updates

$$\forall t, \quad V(s_t) \stackrel{\alpha}{\leftarrow} G_t$$

TD learning

Idea: Online estimation of the **value** function of some policy π
(no need for **terminal** states)

TD updates

$$\forall t, \quad V(s_t) \leftarrow r_t + \gamma V(s_{t+1})$$

cf. Bellman's equation

$$\forall s, \quad V(s) = E_{\pi}(r_t + \gamma V(s_{t+1}) \mid s_t = s)$$

From estimation to control

Let V be the estimate of the value function of some policy π

Policy improvement

$$\pi(s) \leftarrow \arg \max_a \mathbb{E}(r_0 + \gamma V(s_1) | s_0 = s, a_0 = a)$$

Outline

1. Markov decision process
2. Dynamic programming
3. Online estimation
4. **Online control** → SARSA, Q-learning

Action-value function

Recall the **value** function:

$$\forall s, \quad V_{\pi}(s) = E_{\pi}(G | s_0 = s)$$

Without model, how to predict the impact of action a in state s ?

Definition

The **action-value** function is the expected gain from each state-action pair:

$$Q_{\pi}(s, a) = E_{\pi}(G | s_0 = s, a_0 = a)$$

Note that:

$$Q_{\pi}(s, a) = E(r_0 + \gamma V_{\pi}(s_1) | s_0 = s, a_0 = a)$$

Bellman's equation for the action-value function

Proposition

The action-value function Q_π of any policy π is the unique solution to the **fixed-point equation**:

$$\forall s, a, \quad Q(s, a) = E_\pi(r_0 + \gamma Q(s_1, a_1) \mid s_0 = s, a_0 = a)$$

Optimal policy

By definition, the optimal policy is:

$$\pi^*(s) = a^* \in \arg \max Q_*(s, a)$$

with

$$Q_*(s, a) = \mathbb{E}(r_0 + \gamma V_*(s_1) | s_0 = s, a_0 = a)$$

Bellman's optimality theorem

$$\forall s, a, \quad Q_{\pi^*}(s, a) = \max_{\pi} Q_{\pi}(s, a)$$

Bellman's optimality equations

Recall that V_* is the unique solution to:

$$\forall s, \quad V(s) = \max_a E(r_0 + \gamma V(s_1) | s_0 = s, a_0 = a)$$

Proposition

There is a unique solution Q_* to the **fixed-point equation**:

$$\forall s, a, \quad Q(s, a) = E(r_0 + \gamma \max_{a'} Q(s_1, a') | s_0 = s, a_0 = a)$$

Exploration vs exploitation

- Pure exploitation \longrightarrow **greedy**

$$\pi(s) \leftarrow \arg \max_a Q(s, a)$$

- Pure exploration \longrightarrow **random**

$$\pi(s) \leftarrow \text{random}$$

- Exploration-exploitation trade-off \longrightarrow e.g., ϵ -**greedy**

$$\pi(s) \leftarrow \begin{cases} a^* \in \arg \max_a Q(s, a) & \text{with probability } 1 - \epsilon \\ \text{random} & \text{with probability } \epsilon \end{cases}$$

ε -greedy policy improvement

Let π be an ε -greedy policy and:

$$\pi'(s) \leftarrow \begin{cases} a^* \in \arg \max_a Q_\pi(s, a) & \text{with probability } 1 - \varepsilon' \\ \text{random} & \text{with probability } \varepsilon' \end{cases}$$

Proposition

If $\varepsilon' \leq \varepsilon$, the policy π' is better than π :

$$\forall s, \quad V_{\pi'}(s) \geq V_\pi(s)$$

Estimation of the action-value function

MC updates

Given some sequence $s_0, a_0, s_1, a_1, \dots, s_T$,

$$\forall t, \quad Q(s_t, a_t) \stackrel{\alpha}{\leftarrow} G_t$$

TD updates

$$\forall t, \quad Q(s_t, a_t) \stackrel{\alpha}{\leftarrow} r_t + \gamma Q(s_{t+1}, a_{t+1})$$

cf. Bellman's equation

$$\forall s, a, \quad Q(s, a) = \mathbb{E}(r_t + \gamma Q(s_{t+1}, a_{t+1}) \mid s_t = s, a_t = a)$$

SARSA

Key idea: Learn the **optimal** action-value function Q_* online, using the ε -greedy policy:

$$\pi(s) \leftarrow \begin{cases} a^* \in \arg \max_a Q(s, a) & \text{with probability } 1 - \varepsilon \\ \text{random} & \text{with probability } \varepsilon \end{cases}$$

TD updates

$$\forall t, \quad Q(s_t, a_t) \stackrel{\alpha}{\leftarrow} r_t + \gamma Q(s_{t+1}, a_{t+1})$$

Q-learning

Key idea: Learn the **optimal** action-value function Q_* online, using the ε -greedy policy for **control** and the greedy policy for **estimation**.

TD updates


$$\forall t, \quad Q(s_t, a_t) \leftarrow r_t + \gamma \max_a Q(s_{t+1}, a)$$

cf. Bellman's **optimality** equation


$$\forall s, a, \quad Q(s, a) = \mathbb{E}(r_t + \gamma \max_{a'} Q(s_{t+1}, a') \mid s_t = s, a_t = a)$$

Toy example: Pac-Man ($n = 1000$)

SARSA




Q-learning




Toy example: Pac-Man ($n = 100,000$)

SARSA



Q-learning



Outline

1. Markov decision process
2. Dynamic programming → Policy Iteration, Value Iteration
3. Online estimation → Monte-Carlo, TD learning
4. Online control → SARSA, Q-learning

References

Olivier Sigaud

Course on [Reinforcement Learning](#) (slides, videos)

David Silver

Course on [Reinforcement Learning](#) (slides)

Richard Sutton and Andrew Barto

Book [Reinforcement Learning: An Introduction](#) (2015)