



# Deep learning for natural language processing

Chloé Clavel,  
[chloe.clavel@telecom-paristech.fr](mailto:chloe.clavel@telecom-paristech.fr),

Telecom ParisTech, France



# Outline of the course

## Introduction

- Objectives of the course

- Problem statement

## Classical machine learning vs. deep learning

- First option : Classical machine learning

- Second option : Deep learning

- Use for NLP

## NLP Inputs of Neural Networks

- Use pre-trained word vectors

- Train your own word vectors

- Re-train vectors for your task

## Other NN architectures

- Convolutional Neural Networks

- Recurrent neural networks

## Introduction

Classical machine learning vs. deep learning  
NLP Inputs of Neural Networks  
Other NN architectures

Objectives of the course  
Problem statement

# Introduction



## Objectives of the course

At the end of this lecture,

- ▶ you will be able to explain the "philosophy" of deep learning vs. classical machine learning approaches
- ▶ you will master the ML NN architectures for NLP tasks
- ▶ you will be able to cite other neural network architectures for NLP tasks and explain their underlying principles



## Problem statement

- ▶ Training dataset consisting of samples  $\{x_i, y_i\}_{i=1, N}$
- ▶  $x_i$  - inputs, e.g. words (indices or vectors!), context windows, sentences, documents, etc.
- ▶  $y_i$  - labels we try to predict, e.g. other words, class : sentiment, named entities, buy/sell decision,



## NLP tasks

Assigning labels to words :

- ▶ Part-Of-Speech tagging (POS),
- ▶ chunking (CHUNK),
- ▶ Named Entity Recognition (NER)
- ▶ Semantic Role Labeling (SRL)

Assigning labels to sentence/document :

- ▶ Topic classification
- ▶ opinion classification (positive vs. negative)

## Classical machine learning vs. deep learning

# Classical machine learning vs. deep learning

Could speech and language processing be seen as a linear problem ?

## NLP requirements

Input-output functions should solve the selectivity-invariance dilemma

- ▶ insensitive to irrelevant variations of the inputs
- ▶ very sensitive to particular minute variations of the inputs
- ▶ (for example : the pitch variation due to the speaker when you want to develop an emotion recognition system)



## First option : Classical machine learning

### In the simplest cases :

- ▶ linear classifiers on top of **hand-engineered features**
- ▶ A two-class linear classifier computes a weighted sum of the feature vector component
- ▶ if the weighted sum is above a threshold → choose the class

## First option : Classical machine learning

With this option, the challenge is on the design of hand-engineered features

Using semantics, lexicons, etc. (see Lectures 1 and 2) in order to build feature extractor that solves the selectivity-invariance dilemma : build representations that are

- ▶ selective to the aspects of the text that are important for discrimination
- ▶ invariant to irrelevant aspects

Requires engineering skill and linguistic expertise

## Second option : Deep learning

### Statement

- ▶ do not use linguistic expertise and build general purpose learning procedures to automatically learn representations

### Philosophy

- ▶ input : try to pre-process the features as little as possible
- ▶ use a multilayer neural network (NN) architecture trained in an *end-to-end* fashion.
- ▶ ex : use characters as input

## Second option : Deep learning

### Deep learning architecture

Multilayer stack of simple modules

- ▶ subject to learning
- ▶ that computes non-linear input-output mappings
- ▶ that transforms the inputs to increase both the selectivity and the invariance of the representation

## Second option : Deep learning

For example :

- ▶ A multilayer neural network can distort the input space to make the classes of data linearly separable.
- ▶ with a depth of 5 to 20 non-linear layers, a system can implement extremely intricate functions of its inputs that are simultaneously sensitive to minutes details and insensitive to large irrelevant variations
- ▶ If the weights are set correctly, a neural network with enough neurons and a non-linear activation function can approximate a very wide range of mathematical functions

## ML NN use for NLP

- ▶ For binary classification problems
- ▶ For multiclass classification problems
- ▶ More complex structured prediction problems

**Advantages :** The non-linearity of the network, as well as the ability to easily integrate pre-trained word embeddings, often lead to superior classification accuracy.

## ML NN use for NLP

Examples :

- ▶ Syntactic parsing : Chen, D., & Manning, C. (2014). A Fast and Accurate Dependency Parser using Neural Networks. EMNLP 2014
- ▶ Dialog state tracking : Henderson, M., Thomson, B., & Young, S. (2013). Deep Neural Network Approach for the Dialog State Tracking Challenge. Sigdial 2013

## NLP Inputs of Neural Networks



## ML NN Inputs for classification

### Reminder from Lecture 2b about word embeddings

INPUT : one-hot vectors, words are represented as indices taken from a finite dictionary  $\mathcal{D}$

OUTPUT : *Lookup table* dense feature vector  $W = L$

$$L = d \begin{matrix} & |V| \\ \begin{bmatrix} \bullet & \bullet & \dots & \bullet & \dots & \bullet & \bullet \\ \bullet & \bullet & \dots & \bullet & \dots & \bullet & \bullet \\ \bullet & \bullet & \dots & \bullet & \dots & \bullet & \bullet \end{bmatrix} \\ \text{aardvark} \quad \text{a} \quad \dots \quad \text{meta} \quad \dots \quad \text{zebra} \end{matrix}$$

To get the word vector corresponding to the one-hot vector  $e$  :

$$L * e$$

## Use pre-trained word vectors

- ▶ The best option if you have a small training dataset
- ▶ Example : pre-trained word embedding learned on big database, but not specific neither to the data nor to the task (sentiment analysis)
  - In Valentin Barriere, Chloé Clavel, Slim ESSID : « Attitude Classification in Adjacency Pairs of a Human-Agent Interaction with Hidden Conditional Random Fields », ICASSP 2018
  - we had about 500 utterances and we use representations learnt from a Google News corpus of 100 billions words <https://code.google.com/archive/p/word2vec/>

## Train your word vectors on your database in an unsupervised manner

- ▶ The best option if you have a big unlabelled dataset with peculiarities
- ▶ Example :
  - In Maslowski, I., Lagarde, D., Clavel, C., In-the-wild chatbot corpus from opinion analysis to interaction problem detection, ICNLSSP 2017
  - we trained word2vec on 1,813,934 dialogues.

## Re-train vectors for your task

- ▶ The best option when using massive data with a small labelled subset
- ▶ Approach :
  - pretrain a neural encoder or a word2vec on the massive unlabelled data with unsupervised objectives
  - Fine-tune the pretrained model on a specific downstream task using the small labelled dataset

## Re-train vectors for your task

Issue : Pre-trained vectors that do not appear in the training set used for fine-tuning do not change<sup>1</sup>




Figure – Option 1 use of Pretrained word2vec for sentiment analysis




Figure – Option 3 Fine-tuning for sentiment analysis

## Re-train vectors for your task

### Example :

SSWE - Create a word embedding for the task of sentiment analysis with different training objectives to integrate the sentiment information of tweets.

SSWETang, Duyu, et al. "Learning Sentiment-Specific Word Embedding for Twitter Sentiment Classification." ACL (1). 2014.

## SSWE - Training objectives :




Figure 1: The traditional C&W model and our neural networks (SSWE<sub>h</sub> and SSWE<sub>u</sub>) for learning sentiment-specific word embedding.

- ▶ Syntactic : Optimize the prediction of language model score (probability of the n-gram)
- ▶ Sentiment : Optimize the prediction of sentiment score of the n-gram so that it should be consistent with the gold polarity annotation of sentence

## Re-train vectors for your task

Another Example : Johnson, R., & Zhang, T. (2015).  
Semi-supervised convolutional neural networks for text  
categorization via region embedding. In Advances in neural  
information processing systems



## Train multilayer neural network (NN) architecture in an end-to-end fashion

The best option if you have a sufficiently big **labelled** dataset

STEP 1 : The architecture takes the input sentences and learns several layers of feature extraction that process the inputs.

STEP 2 (fine-tuning) : The features computed by the deep layers of the network are automatically trained by backpropagation to be relevant to the task.


## Other NN architectures

## Other NN architectures

- ▶ Convolutional neural networks
- ▶ Recurrent neural networks and variants

# Convolutional Neural Networks

- ▶ Variation of multilayer perceptrons designed to require minimal preprocessing and using *convolutional* layers
- ▶ The network learns the filters



# Convolutional Neural Networks

Example of use for the text : Johnson, R., & Zhang, T. (2014).  
Effective use of word order for text categorization with  
convolutional neural networks.





Figure 3: Convolution layer for variable-sized text.


- + CNN can model semantic clues in contextual windows
- CNN have difficulty to preserve sequential orders and to model long-term contextual information

# Recurrent Neural Networks



- Read inputs  $x_t$  to accumulate state  $h_t$  and predict outputs  $y_t$ .  
ex : Use for language models (output = next word)

# Recurrent Neural Networks




- $h_t$  contains information about the whole past sequence

$$h_t = f(h_{t-1}, x_t, \theta)$$

- the network learns to use  $h_t$  as a kind of lossy summary of the task-relevant aspects of the past sequence of inputs up to  $t$

# Recurrent Neural Networks




- ▶ autoencoder framework : we ask  $h_t$  to be rich enough to allow one to approximately recover the input sequence, as in autoencoder framework
- ▶ Variants : LSTM networks (Long Short Term Memory Networks), RNN using gating mechanisms such as GRU (Gated Recurrent Units)



## Recursive tensor network


Another generalization of recurrent networks with deep tree structure rather than the chain-like structure of RNN



R. Socher, A. Perelygin, J. Wu, J. Chuang, C. D. Manning, A. Y. Ng, and C. Potts, Recursive deep models for semantic compositionality over a sentiment treebank, EMNLP 2013


## Recursive deep models

- ▶ Database : Treebank Sentiment sentences of movie reviews parsed with the Stanford parser and represented by a tree
- ▶ Each node of the tree is labelled in  $(-, +, 0)$  to provide the structure that is required for the training of a recursive model



## Recursive deep models

- ▶ Training step : learning  $g$  function that compute the upper outputs in the binary tree
- ▶ Decision step : recursively apply the activation functions :



## Support and materials

- ▶ LeCun, Yann, Yoshua Bengio, and Geoffrey Hinton. "Deep learning." *nature* 521.7553 (2015) : 436.
- ▶ Lectures from Stanford  
<http://cs224d.stanford.edu/lectures/CS224d-Lecture4.pdf>
- ▶ Collobert, R., Weston, J., Bottou, L., Karlen, M., Kavukcuoglu, K., & Kuksa, P. (2011). Natural language processing (almost) from scratch. *Journal of Machine Learning Research*, 12(Aug), 2493-2537.
- ▶ Goldberg, Yoav. "A primer on neural network models for natural language processing." *Journal of Artificial Intelligence Research* 57 (2016) : 345-420.
- ▶ Lectures from Oxford :  
<https://github.com/oxford-cs-deepnlp-2017/lectures>